

EstiNet MT198T OpenFlow Command Reference

USER GUIDE

Table of Contents

1	General Guide	1
2	ovs-appctl	1
2.1	COMMON COMMANDS	2
2.1.1	GENERAL COMMANDS	2
2.1.2	LOGGING COMMANDS	2
2.2	OPTIONS	5
3	ovs-ofctl	6
3.1	OpenFlow Switch Management Commands	6
3.1.1	show switch	6
3.1.2	dump-tables switch	7
3.1.3	dump-table-features switch	7
3.1.4	dump-ports switch [netdev]	7
3.1.5	dump-ports-desc switch [port]	7
3.1.6	mod-port switch port action	7
3.1.7	get-frags switch	8
3.1.8	dump-flows switch [flows]	8
3.1.9	dump-aggregate switch [flows]	8
3.2	OpenFlow Group Table Commands	8
3.2.1	dump-groups switch [group]	9
3.2.2	dump-group-features switch	9
3.2.3	dump-group-stats switch [groups]	9
3.3	OpenFlow Switch Meter Table Commands	9
3.3.1	add-meter switch meter	9
3.3.2	mod-meter switch meter	9
3.3.3	del-meters switch	9
3.3.4	del-meter switch [meter]	9
3.3.5	dump-meters switch	9
3.3.6	dump-meter switch [meter]	9
3.3.7	meter-stats switch [meter]	9
3.3.8	meter-features switch	9
3.4	OpenFlow Switch Flow Table Commands	10
3.5	OpenFlow Switch Group Table Commands	11
3.6	OpenFlow Switch Monitoring Commands	11
3.6.1	snoop switch	11
3.6.2	monitor switch [miss-len] [invalid_ttl] [watch:[spec...]]	12
3.7	OpenFlow Switch and Controller Commands	13
3.7.1	probe target	13
3.7.2	ping target [n]	13
3.7.3	benchmark target n count	13
3.8	Other Commands	13
3.8.1	ofp-parse file	13
3.8.2	ofp-parse-pcap file [port...]	14
3.9	Flow Syntax	14
3.10	Group Syntax	46
3.11	Meter Syntax	49
3.12	OPTIONS	50
4	ovs-vsctl	58
4.1	Open vSwitch Commands	64
4.2	Bridge Commands	65
4.3	Port Commands	66
4.4	Interface Commands	67
4.5	OpenFlow Controller Connectivity	67
4.6	Controller Failure Settings	68
4.7	Manager Connectivity	69
4.8	SSL Configuration	70
4.9	Auto-Attach Commands	71
4.10	Database Commands	71

1 General Guide

EstiNet MT198T switch supports most of the OpenFlow protocol version 1.3.4 functions . It is installed with the Open vSwitch agent, which is responsible for handling OpenFlow commands from an SDN controller or a local CLI console. This way, users can use the OpenFlow functions provided by the switch via OVS commands, which are explained in detail in following chapters.

2 ovs-appctl

NAME

ovs-appctl - utility for configuring running Open vSwitch daemons

SYNOPSIS

```
ovs-appctl [--target=target | -t target] command [arg...]  
ovs-appctl --help  
ovs-appctl --version
```

DESCRIPTION

Open vSwitch daemons accept certain commands at runtime to control their behavior and query their settings. Every daemon accepts a common set of commands documented under COMMON COMMANDS below. Some daemons support additional commands documented in their own manpages.

ovs-vswitchd in particular accepts a number of additional commands documented in ovs-vswitchd(8).

The ovs-appctl program provides a simple way to invoke these commands.

The command to be sent is specified on ovs-appctl's command line as non-option arguments. ovs-appctl sends the command and prints the daemon's response on standard output.

In normal use only a single option is accepted:

```
-t target  
--target=target  
    Tells ovs-appctl which daemon to contact.
```

If target begins with / it must name a Unix domain socket on which an Open vSwitch daemon is listening for control channel connections. By default, each daemon listens on a Unix domain socket named `//var/run/openvswitch/program.pidctl`, where program is the program's name and pid is its process ID. For example, if ovs-vswitchd has PID 123, it would listen on `//var/run/openvswitch/ovs-vswitchd.123ctl`.

Otherwise, ovs-appctl looks for a pidfile, that is, a file whose contents are the process ID of a running process as a decimal number, named `//var/run/openvswitch/target.pid`. (The `--pidfile` option makes an Open vSwitch daemon create a pidfile.) ovs-appctl reads the pidfile, then looks for a Unix socket named `//var/run/openvswitch/target.pidctl`, where pid is replaced by the process ID read from the pidfile, and uses that file as if it had been specified directly as the target.

On Windows, target can be an absolute path to a file that contains a localhost TCP port on which an Open vSwitch daemon is listening for control channel connections. By default, each daemon writes the TCP port on which it is listening for control

connection into the file `program.ctl` located inside the configured `OVS_RUNDIR` directory. If `target` is not an absolute path, `ovs-appctl` looks for a file named `target.ctl` in the configured `OVS_RUNDIR` directory.

The default target is `ovs-vswitchd`.

2.1 COMMON COMMANDS

Every Open vSwitch daemon supports a common set of commands, which are documented in this section.

2.1.1 GENERAL COMMANDS

These commands display daemon-specific commands and the running version. Note that these commands are different from the `--help` and `--version` options that return information about the `ovs-appctl` utility itself.

`list-commands`

Lists the commands supported by the target.

`version`

Displays the version and compilation date of the target.

2.1.2 LOGGING COMMANDS

Open vSwitch has several log levels. The highest-severity log level is:

`off` No message is ever logged at this level, so setting a logging destination's log level to `off` disables logging to that destination.

The following log levels, in order of descending severity, are available:

`emer` A major failure forced a process to abort.

`err` A high-level operation or a subsystem failed. Attention is warranted.

`warn` A low-level operation failed, but higher-level subsystems may be able to recover.

`info` Information that may be useful in retrospect when investigating a problem.

`dbg` Information useful only to someone with intricate knowledge of the system, or that would commonly cause too-voluminous log output. Log messages at this level are not logged by default.

Every Open vSwitch daemon supports the following commands for examining and adjusting log levels.

`vlog/list`

Lists the known logging modules and their current levels.

`vlog/list-pattern`

Lists logging pattern used for each destination.

`vlog/set [spec]`

Sets logging levels. Without any spec, sets the log level for every module and destination to dbg. Otherwise, spec is a list of words separated by spaces or commas or colons, up to one from each category below:

- A valid module name, as displayed by the `vlog/list` command on `ovs-appctl(8)`, limits the log level change to the specified module.
- `syslog`, `console`, or `file`, to limit the log level change to only to the system log, to the console, or to a file, respectively.

On Windows platform, `syslog` is accepted as a word and is only useful if the target was started with the `--syslog-target` option (the word has no effect otherwise).

- `off`, `emer`, `err`, `warn`, `info`, or `dbg`, to control the log level. Messages of the given severity or higher will be logged, and messages of lower severity will be filtered out. `off` filters out all messages.

Case is not significant within spec.

Regardless of the log levels set for file, logging to a file will not take place unless the target application was invoked with the `--log-file` option.

For compatibility with older versions of OVS, any is accepted as a word but has no effect.

`vlog/set PATTERN:destination:pattern`

Sets the log pattern for destination to pattern. Each time a message is logged to destination, pattern determines the message's formatting. Most characters in pattern are copied literally to the log, but special escapes beginning with % are expanded as follows:

%A The name of the application logging the message, e.g. `ovs-vswitchd`.

%B The RFC5424 syslog PRI of the message.

%C The name of the module (as shown by `ovs-appctl --list`) logging the message.

%d The current date and time in ISO 8601 format (YYYY-MM-DD HH:MM:SS).

%d{format}

The current date and time in the specified format, which takes the same format as the template argument to `strftime(3)`. As an extension, any # characters in format will be replaced by fractional seconds, e.g. use `%H:%M:%S.###` for the time to the nearest millisecond. Sub-second times are only approximate and currently decimal places after the third will always be reported as zero.

%D The current UTC date and time in ISO 8601 format (YYYY-MM-DD HH:MM:SS).

`%D{format}` The current UTC date and time in the specified format, which takes the same format as the template argument to `strftime(3)`. Supports the same extension for sub-second resolution as `%d{...}`.

`%E` The hostname of the node running the application.

`%m` The message being logged.

`%N` A serial number for this message within this run of the program, as a decimal number. The first message a program logs has serial number 1, the second one has serial number 2, and so on.

`%n` A new-line.

`%p` The level at which the message is logged, e.g. `DBG`.

`%P` The program's process ID (`pid`), as a decimal number.

`%r` The number of milliseconds elapsed from the start of the application to the time the message was logged.

`%t` The subprogram name, that is, an identifying name for the process or thread that emitted the log message, such as `monitor` for the process used for `--monitor` or `main` for the primary process or thread in a program.

`%T` The subprogram name enclosed in parentheses, e.g. `(monitor)`, or the empty string for the primary process or thread in a program.

`%%` A literal `%`.

A few options may appear between the `%` and the format specifier character, in this order:

- Left justify the escape's expansion within its field width. Right justification is the default.
- 0 Pad the field to the field width with 0s. Padding with spaces is the default.
- width A number specifies the minimum field width. If the escape expands to fewer characters than width then it is padded to fill the field width. (A field wider than width is not truncated to fit.)

The default pattern for console and file output is `%D{%Y-%m-%dT%H:%M:%SZ}|%05N|%c|%p|%m`; for syslog output, `%05N|%c|%p|%m`.

Daemons written in Python (e.g. `ovs-xapi-sync`, `ovs-monitor-ipsec`) do not allow control over the log pattern.

`vlog/set FACILITY:facility`

Sets the RFC5424 facility of the log message. facility can be one of `kern`, `user`, `mail`, `daemon`, `auth`, `syslog`, `lpr`, `news`, `uucp`, `clock`, `ftp`, `ntp`, `audit`, `alert`, `clock2`, `local0`, `local1`, `local2`, `local3`, `local4`, `local5`, `local6` or `local7`.

vlog/close

Causes the daemon to close its log file, if it is open. (Use vlog/reopen to reopen it later.)

vlog/reopen

Causes the daemon to close its log file, if it is open, and then reopen it. (This is useful after rotating log files, to cause a new log file to be used.)

This has no effect if the target application was not invoked with the --log-file option.

2.2 OPTIONS

-h

--help Prints a brief help message to the console.

-V

--version

Prints version information to the console.

3 ovs-ofctl

NAME

ovs-ofctl - administer OpenFlow switches

SYNOPSIS

ovs-ofctl [options] command [switch] [args...]

DESCRIPTION

The ovs-ofctl program is a command line tool for monitoring and administering OpenFlow switches. It can also show the current state of an OpenFlow switch, including features, configuration, and table entries. It should work with any OpenFlow switch, not just Open vSwitch.

3.1 OpenFlow Switch Management Commands

These commands allow ovs-ofctl to monitor and administer an OpenFlow switch. It is able to show the current state of a switch, including features, configuration, and table entries.

Most of these commands take an argument that specifies the method for connecting to an OpenFlow switch. The following connection methods are supported:

ssl:ip[:port]

tcp:ip[:port]

The specified port on the host at the given ip, which must be expressed as an IP address (not a DNS name) in IPv4 or IPv6 address format. Wrap IPv6 addresses in square brackets, e.g. tcp:[::1]:6653. For ssl, the --private-key, --certificate, and --ca-cert options are mandatory.

If port is not specified, it defaults to 6653.

unix:file

On POSIX, a Unix domain server socket named file.

On Windows, a localhost TCP port written in file.

file This is short for unix:file, as long as file does not contain a colon.

bridge This is short for unix://var/run/openvswitch/bridge.mgmt, as long as bridge does not contain a colon.

[type@]dp

Attempts to look up the bridge associated with dp and open as above. If type is given, it specifies the datapath provider of dp, otherwise the default provider system is assumed.

3.1.1 show switch

Prints to the console information on switch, including information on its flow tables and ports.

3.1.2 `dump-tables switch`

Prints to the console statistics for each of the flow tables used by switch.

3.1.3 `dump-table-features switch`

Prints to the console features for each of the flow tables used by switch.

3.1.4 `dump-ports switch [netdev]`

Prints to the console statistics for network devices associated with switch. If netdev is specified, only the statistics associated with that device will be printed. netdev can be an OpenFlow assigned port number or device name, e.g. eth0.

3.1.5 `dump-ports-desc switch [port]`

Prints to the console detailed information about network devices associated with switch. To dump only a specific port, specify its number as port. Otherwise, if port is omitted, or if it is specified as ANY, then all ports are printed. This is a subset of the information provided by the show command.

If the connection to switch negotiates OpenFlow 1.0, 1.2, or 1.2, this command uses an OpenFlow extension only implemented in Open vSwitch (version 1.7 and later).

Only OpenFlow 1.5 and later support dumping a specific port. Earlier versions of OpenFlow always dump all ports.

3.1.6 `mod-port switch port action`

Modify characteristics of port port in switch. port may be an OpenFlow port number or name or the keyword LOCAL (the preferred way to refer to the OpenFlow local port). The action may be any one of the following:

up

down Enable or disable the interface. This is equivalent to ifconfig up or ifconfig down on a Unix system.

stp

no-stp Enable or disable 802.1D spanning tree protocol (STP) on the interface. OpenFlow implementations that don't support STP will refuse to enable it.

receive

no-receive

receive-stp

no-receive-stp

Enable or disable OpenFlow processing of packets received on this interface. When packet processing is disabled, packets will be dropped instead of being processed through the OpenFlow table. The receive or no-receive setting applies to all packets except 802.1D spanning tree packets, which are separately controlled by receive-stp or no-receive-stp.

forward

no-forward

Allow or disallow forwarding of traffic to this interface. By default, forwarding is enabled.

flood
no-flood

Controls whether an OpenFlow flood action will send traffic out this interface. By default, flooding is enabled. Disabling flooding is primarily useful to prevent loops when a spanning tree protocol is not in use.

packet-in
no-packet-in

Controls whether packets received on this interface that do not match a flow table entry generate a ``packet in'' message to the OpenFlow controller. By default, ``packet in'' messages are enabled.

The show command displays (among other information) the configuration that mod-port changes.

3.1.7 get-frags switch

Prints switch's fragment handling mode. See set-frags, below, for a description of each fragment handling mode.

The show command also prints the fragment handling mode among its other output.

3.1.8 dump-flows switch [flows]

Prints to the console all flow entries in switch's tables that match flows. If flows is omitted, all flows in the switch are retrieved. See Flow Syntax, below, for the syntax of flows. The output format is described in Table Entry Output.

By default, ovs-ofctl prints flow entries in the same order that the switch sends them, which is unlikely to be intuitive or consistent. See the description of --sort and --rsort, under OPTIONS below, to influence the display order.

3.1.9 dump-aggregate switch [flows]

Prints to the console aggregate statistics for flows in switch's tables that match flows. If flows is omitted, the statistics are aggregated across all flows in the switch's flow tables. See Flow Syntax, below, for the syntax of flows. The output format is described in Table Entry Output.

3.2 OpenFlow Group Table Commands

The following commands work only with switches that support OpenFlow 1.1 or later. Because support for OpenFlow 1.1 and later is still experimental in Open vSwitch, it is necessary to explicitly enable these protocol versions in ovs-ofctl (using -O) and in the switch itself (with the protocols column in the Bridge table). For more information, see ``Q: What versions of OpenFlow does Open vSwitch support?'' in the Open vSwitch FAQ.

3.2.1 **dump-groups switch [group]**

Prints group entries in switch's tables to console. To dump only a specific group, specify its number as group. Otherwise, if group is omitted, or if it is specified as ALL, then all groups are printed. Each line of output is a group entry as described in Group Syntax below.

Only OpenFlow 1.5 and later support dumping a specific group. Earlier versions of OpenFlow always dump all groups.

3.2.2 **dump-group-features switch**

Prints to the console the group features of the switch.

3.2.3 **dump-group-stats switch [groups]**

Prints to the console statistics for the specified groups in the switch's tables. If groups is omitted then statistics for all groups are printed. See Group Syntax, below, for the syntax of groups.

3.3 **OpenFlow Switch Meter Table Commands**

These commands manage the meter table in an OpenFlow switch. In each case, meter specifies a meter entry in the format described in Meter Syntax, below.

OpenFlow 1.3 introduced support for meters, so these commands only work with switches that support OpenFlow 1.3 or later. The caveats described for groups in the previous section also apply to meters.

3.3.1 **add-meter switch meter**

Add a meter entry to switch's tables. The meter syntax is described in section Meter Syntax, below.

3.3.2 **mod-meter switch meter**

Modify an existing meter.

3.3.3 **del-meters switch**

3.3.4 **del-meter switch [meter]**

Delete entries from switch's meter table. meter can specify a single meter with syntax meter=id, or all meters with syntax meter=all.

3.3.5 **dump-meters switch**

3.3.6 **dump-meter switch [meter]**

Print meter configuration. meter can specify a single meter with syntax meter=id, or all meters with syntax meter=all.

3.3.7 **meter-stats switch [meter]**

Print meter statistics. meter can specify a single meter with syntax meter=id, or all meters with syntax meter=all.

3.3.8 **meter-features switch**

Print meter features.

3.4 OpenFlow Switch Flow Table Commands

These commands manage the flow table in an OpenFlow switch. In each case, flow specifies a flow entry in the format described in Flow Syntax, below, file is a text file that contains zero or more flows in the same syntax, one per line, and the optional --bundle option operates the command as a single atomic transaction, see option --bundle, below.

```
[--bundle] add-flow switch flow
[--bundle] add-flow switch - < file
[--bundle] add-flows switch file
```

Add each flow entry to switch's tables. Each flow specification (e.g., each line in file) may start with add, modify, delete, modify_strict, or delete_strict keyword to specify whether a flow is to be added, modified, or deleted, and whether the modify or delete is strict or not. For backwards compatibility a flow specification without one of these keywords is treated as a flow add. All flow mods are executed in the order specified.

```
[--bundle] [--strict] mod-flows switch flow
[--bundle] [--strict] mod-flows switch - < file
```

Modify the actions in entries from switch's tables that match the specified flows. With --strict, wildcards are not treated as active for matching purposes.

```
[--bundle] del-flows switch
[--bundle] [--strict] del-flows switch [flow]
[--bundle] [--strict] del-flows switch - < file
```

Deletes entries from switch's flow table. With only a switch argument, deletes all flows. Otherwise, deletes flow entries that match the specified flows. With --strict, wildcards are not treated as active for matching purposes.

```
[--bundle] [--readd] replace-flows switch file
```

Reads flow entries from file (or stdin if file is -) and queries the flow table from switch. Then it fixes up any differences, adding flows from flow that are missing on switch, deleting flows from switch that are not in file, and updating flows in switch whose actions, cookie, or timeouts differ in file.

With --readd, ovs-ofctl adds all the flows from file, even those that exist with the same actions, cookie, and timeout in switch. This resets all the flow packet and byte counters to 0, which can be useful for debugging.

```
diff-flows source1 source2
```

Reads flow entries from source1 and source2 and prints the differences. A flow that is in source1 but not in source2 is printed preceded by a -, and a flow that is in source2 but not in source1 is printed preceded by a +. If a flow exists in both source1 and source2 with different actions, cookie, or timeouts, then both versions are printed preceded by - and +, respectively.

source1 and source2 may each name a file or a switch. If a name begins with / or ., then it is considered to be a file name. A name that contains : is considered to be a switch. Otherwise, it is a file if a file by that name exists, a switch if not.

For this command, an exit status of 0 means that no differences

were found, 1 means that an error occurred, and 2 means that some differences were found.

packet-out switch in_port actions packet...

Connects to switch and instructs it to execute the OpenFlow actions on each packet. Each packet is specified as a series of hex digits. For the purpose of executing the actions, the packets are considered to have arrived on in_port, which may be an OpenFlow port number or name (e.g. eth0), the keyword LOCAL (the preferred way to refer to the OpenFlow 'local' port), or the keyword NONE to indicate that the packet was generated by the switch itself.

3.5 OpenFlow Switch Group Table Commands

These commands manage the group table in an OpenFlow switch. In each case, group specifies a group entry in the format described in Group Syntax, below, and file is a text file that contains zero or more groups in the same syntax, one per line.

```
add-group switch group
add-group switch - < file
add-groups switch file
```

Add each group entry to switch's tables.

```
[--may-create] mod-group switch group
[--may-create] mod-group switch - < file
```

Modify the action buckets in entries from switch's tables for each group entry. If a specified group does not already exist, then without --may-create, this command has no effect; with --may-create, it creates a new group. The --may-create option uses an Open vSwitch extension to OpenFlow only implemented in Open vSwitch 2.6 and later.

```
del-groups switch
del-groups switch [group]
del-groups switch - < file
```

Deletes entries from switch's group table. With only a switch argument, deletes all groups. Otherwise, deletes the group for each group entry.

```
insert-buckets switch group
insert-buckets switch - < file
```

Add buckets to an existing group present in the switch's group table. If no command_bucket_id is present in the group specification then all buckets of the group are removed.

```
remove-buckets switch group
remove-buckets switch - < file
```

Remove buckets to an existing group present in the switch's group table. If no command_bucket_id is present in the group specification then all buckets of the group are removed.

3.6 OpenFlow Switch Monitoring Commands

3.6.1 snoop switch

Connects to switch and prints to the console all OpenFlow messages received. Unlike other ovs-ofctl commands, if switch is the name of a bridge, then the snoop command connects to a Unix

domain socket named `//var/run/openvswitch/switch.snoop`. `ovs-vswitchd` listens on such a socket for each bridge and sends to it all of the OpenFlow messages sent to or received from its configured OpenFlow controller. Thus, this command can be used to view OpenFlow protocol activity between a switch and its controller.

When a switch has more than one controller configured, only the traffic to and from a single controller is output. If none of the controllers is configured as a master or a slave (using a Nicira extension to OpenFlow 1.0 or 1.1, or a standard request in OpenFlow 1.2 or later), then a controller is chosen arbitrarily among them. If there is a master controller, it is chosen; otherwise, if there are any controllers that are not masters or slaves, one is chosen arbitrarily; otherwise, a slave controller is chosen arbitrarily. This choice is made once at connection time and does not change as controllers reconfigure their roles.

If a switch has no controller configured, or if the configured controller is disconnected, no traffic is sent, so monitoring will not show any traffic.

3.6.2 `monitor switch [miss-len] [invalid_ttl] [watch:[spec...]]`

Connects to switch and prints to the console all OpenFlow messages received. Usually, switch should specify the name of a bridge in the `ovs-vswitchd` database.

If `miss-len` is provided, `ovs-ofctl` sends an OpenFlow ```set configuration``` message at connection setup time that requests `miss-len` bytes of each packet that misses the flow table. Open vSwitch does not send these and other asynchronous messages to an `ovs-ofctl` monitor client connection unless a nonzero value is specified on this argument. (Thus, if `miss-len` is not specified, very little traffic will ordinarily be printed.)

If `invalid_ttl` is passed, `ovs-ofctl` sends an OpenFlow ```set configuration``` message at connection setup time that requests `INVALID_TTL_TO_CONTROLLER`, so that `ovs-ofctl` monitor can receive ```packet-in``` messages when TTL reaches zero on `dec_ttl` action. Only OpenFlow 1.1 and 1.2 support `invalid_ttl`; Open vSwitch also implements it for OpenFlow 1.0 as an extension.

`watch:[spec...]` causes `ovs-ofctl` to send a ```monitor request``` Nicira extension message to the switch at connection setup time. This message causes the switch to send information about flow table changes as they occur. The following comma-separated spec syntax is available:

```
!initial
    Do not report the switch's initial flow table contents.

!add
    Do not report newly added flows.

!delete
    Do not report deleted flows.

!modify
```

Do not report modifications to existing flows.

`!own` Abbreviate changes made to the flow table by `ovs-ofctl`'s own connection to the switch. (These could only occur using the `ofctl/send` command described below under RUN-TIME MANAGEMENT COMMANDS.)

`!actions`
Do not report actions as part of flow updates.

`table=number`
Limits the monitoring to the table with the given number between 0 and 254. By default, all tables are monitored.

`out_port=port`
If set, only flows that output to port are monitored. The port may be an OpenFlow port number or keyword (e.g. LOCAL).

`field=value`
Monitors only flows that have field specified as the given value. Any syntax valid for matching on dump-flows may be used.

This command may be useful for debugging switch or controller implementations. With `watch:`, it is particularly useful for observing how a controller updates flow tables.

3.7 OpenFlow Switch and Controller Commands

The following commands, like those in the previous section, may be applied to OpenFlow switches, using any of the connection methods described in that section. Unlike those commands, these may also be applied to OpenFlow controllers.

3.7.1 probe target

Sends a single OpenFlow echo-request message to target and waits for the response. With the `-t` or `--timeout` option, this command can test whether an OpenFlow switch or controller is up and running.

3.7.2 ping target [n]

Sends a series of 10 echo request packets to target and times each reply. The echo request packets consist of an OpenFlow header plus `n` bytes (default: 64) of randomly generated payload. This measures the latency of individual requests.

3.7.3 benchmark target n count

Sends `count` echo request packets that each consist of an OpenFlow header plus `n` bytes of payload and waits for each response. Reports the total time required. This is a measure of the maximum bandwidth to target for round-trips of `n`-byte messages.

3.8 Other Commands

3.8.1 ofp-parse file

Reads file (or stdin if file is `-`) as a series of OpenFlow mes-

sages in the binary format used on an OpenFlow connection, and prints them to the console. This can be useful for printing OpenFlow messages captured from a TCP stream.

3.8.2 ofp-parse-pcap file [port...]

Reads file, which must be in the PCAP format used by network capture tools such as tcpdump or Wireshark, extracts all the TCP streams for OpenFlow connections, and prints the OpenFlow messages in those connections in human-readable format on stdout.

OpenFlow connections are distinguished by TCP port number. Non-OpenFlow packets are ignored. By default, data on TCP ports 6633 and 6653 are considered to be OpenFlow. Specify one or more port arguments to override the default.

This command cannot usefully print SSL encrypted traffic. It does not understand IPv6.

3.9 Flow Syntax

Some `ovs-ofctl` commands accept an argument that describes a flow or flows. Such flow descriptions comprise a series field=value assignments, separated by commas or white space. (Embedding spaces into a flow description normally requires quoting to prevent the shell from breaking the description into multiple arguments.)

Flow descriptions should be in normal form. This means that a flow may only specify a value for an L3 field if it also specifies a particular L2 protocol, and that a flow may only specify an L4 field if it also specifies particular L2 and L3 protocol types. For example, if the L2 protocol type `dl_type` is wildcarded, then L3 fields `nw_src`, `nw_dst`, and `nw_proto` must also be wildcarded. Similarly, if `dl_type` or `nw_proto` (the L3 protocol type) is wildcarded, so must be the L4 fields `tcp_dst` and `tcp_src`. `ovs-ofctl` will warn about flows not in normal form.

The following field assignments describe how a flow matches a packet. If any of these assignments is omitted from the flow syntax, the field is treated as a wildcard; thus, if all of them are omitted, the resulting flow matches all packets. The string `*` may be specified to explicitly mark any of these fields as a wildcard. (`*` should be quoted to protect it from shell expansion.)

`in_port=port`

Matches OpenFlow port `port`, which may be an OpenFlow port number or keyword (e.g. `LOCAL`). `ovs-ofctl show`.

(The `resubmit` action can search OpenFlow flow tables with arbitrary `in_port` values, so flows that match port numbers that do not exist from an OpenFlow perspective can still potentially be matched.)

`dl_vlan=vlan`

Matches IEEE 802.1q Virtual LAN tag `vlan`. Specify `0xffff` as `vlan` to match packets that are not tagged with a Virtual LAN; otherwise, specify a number between 0 and 4095, inclusive, as the 12-bit VLAN ID to match.

`dl_vlan_pcp=priority`

Matches IEEE 802.1q Priority Code Point (PCP) priority, which is

specified as a value between 0 and 7, inclusive. A higher value indicates a higher frame priority level.

`dl_src=xx:xx:xx:xx:xx:xx`

`dl_dst=xx:xx:xx:xx:xx:xx`

Matches an Ethernet source (or destination) address specified as 6 pairs of hexadecimal digits delimited by colons (e.g. 00:0A:E4:25:6B:B0).

`dl_src=xx:xx:xx:xx:xx:xx/xx:xx:xx:xx:xx:xx`

`dl_dst=xx:xx:xx:xx:xx:xx/xx:xx:xx:xx:xx:xx`

Matches an Ethernet destination address specified as 6 pairs of hexadecimal digits delimited by colons (e.g. 00:0A:E4:25:6B:B0), with a wildcard mask following the slash. Open vSwitch 1.8 and later support arbitrary masks for source and/or destination. Earlier versions only support masking the destination with the following masks:

`01:00:00:00:00:00`

Match only the multicast bit. Thus, `dl_dst=01:00:00:00:00:00/01:00:00:00:00:00` matches all multicast (including broadcast) Ethernet packets, and `dl_dst=00:00:00:00:00:00/01:00:00:00:00:00` matches all unicast Ethernet packets.

`fe:ff:ff:ff:ff:ff`

Match all bits except the multicast bit. This is probably not useful.

`ff:ff:ff:ff:ff:ff`

Exact match (equivalent to omitting the mask).

`00:00:00:00:00:00`

Wildcard all bits (equivalent to `dl_dst=*`).

`dl_type=ethertype`

Matches Ethernet protocol type `ethertype`, which is specified as an integer between 0 and 65535, inclusive, either in decimal or as a hexadecimal number prefixed by 0x (e.g. 0x0806 to match ARP packets).

`nw_src=ip[/netmask]`

`nw_dst=ip[/netmask]`

When `dl_type` is 0x0800 (possibly via shorthand, e.g. `ip` or `tcp`), matches IPv4 source (or destination) address `ip`, which may be specified as an IP address or host name (e.g. 192.168.1.1 or `www.example.com`). The optional `netmask` allows restricting a match to an IPv4 address prefix. The `netmask` may be specified as a dotted quad (e.g. 192.168.1.0/255.255.255.0) or as a CIDR block (e.g. 192.168.1.0/24). Open vSwitch 1.8 and later support arbitrary dotted quad masks; earlier versions support only CIDR masks, that is, the dotted quads that are equivalent to some CIDR block.

When `dl_type=0x0806` or `arp` is specified, matches the `ar_spa` or `ar_tpa` field, respectively, in ARP packets for IPv4 and Ethernet.

When `dl_type=0x8035` or `rarp` is specified, matches the `ar_spa` or `ar_tpa` field, respectively, in RARP packets for IPv4 and Ethernet.

net.

When `dl_type` is wildcarded or set to a value other than `0x0800`, `0x0806`, or `0x8035`, the values of `nw_src` and `nw_dst` are ignored (see Flow Syntax above).

`nw_proto=proto`

`ip_proto=proto`

When `ip` or `dl_type=0x0800` is specified, matches IP protocol type `proto`, which is specified as a decimal number between 0 and 255, inclusive (e.g. 1 to match ICMP packets or 6 to match TCP packets).

When `ipv6` or `dl_type=0x86dd` is specified, matches IPv6 header type `proto`, which is specified as a decimal number between 0 and 255, inclusive (e.g. 58 to match ICMPv6 packets or 6 to match TCP). The header type is the terminal header as described in the DESIGN document.

When `arp` or `dl_type=0x0806` is specified, matches the lower 8 bits of the ARP opcode. ARP opcodes greater than 255 are treated as 0.

When `rarp` or `dl_type=0x8035` is specified, matches the lower 8 bits of the ARP opcode. ARP opcodes greater than 255 are treated as 0.

When `dl_type` is wildcarded or set to a value other than `0x0800`, `0x0806`, `0x8035` or `0x86dd`, the value of `nw_proto` is ignored (see Flow Syntax above).

`nw_tos=tos`

Matches IP ToS/DSCP or IPv6 traffic class field `tos`, which is specified as a decimal number between 0 and 255, inclusive. Note that the two lower reserved bits are ignored for matching purposes.

When `dl_type` is wildcarded or set to a value other than `0x0800` or `0x86dd`, the value of `nw_tos` is ignored (see Flow Syntax above).

`ip_dscp=dscp`

Matches IP ToS/DSCP or IPv6 traffic class field `dscp`, which is specified as a decimal number between 0 and 63, inclusive.

When `dl_type` is wildcarded or set to a value other than `0x0800` or `0x86dd`, the value of `ip_dscp` is ignored (see Flow Syntax above).

`nw_ecn=ecn`

`ip_ecn=ecn`

Matches ecn bits in IP ToS or IPv6 traffic class fields, which is specified as a decimal number between 0 and 3, inclusive.

When `dl_type` is wildcarded or set to a value other than `0x0800` or `0x86dd`, the value of `nw_ecn` is ignored (see Flow Syntax above).

`nw_ttl=ttl`

Matches IP TTL or IPv6 hop limit value `ttl`, which is specified as a decimal number between 0 and 255, inclusive.

When `dl_type` is wildcarded or set to a value other than `0x0800` or `0x86dd`, the value of `nw_ttl` is ignored (see Flow Syntax above).

```
tcp_src=port
tcp_dst=port
udp_src=port
udp_dst=port
sctp_src=port
sctp_dst=port
```

Matches a TCP, UDP, or SCTP source or destination port `port`, which is specified as a decimal number between 0 and 65535, inclusive.

When `dl_type` and `nw_proto` are wildcarded or set to values that do not indicate an appropriate protocol, the values of these settings are ignored (see Flow Syntax above).

```
tcp_src=port/mask
tcp_dst=port/mask
udp_src=port/mask
udp_dst=port/mask
sctp_src=port/mask
sctp_dst=port/mask
```

Bitwise match on TCP (or UDP or SCTP) source or destination port. The port and mask are 16-bit numbers written in decimal or in hexadecimal prefixed by `0x`. Each 1-bit in mask requires that the corresponding bit in port must match. Each 0-bit in mask causes the corresponding bit to be ignored.

Bitwise matches on transport ports are rarely useful in isolation, but a group of them can be used to reduce the number of flows required to match on a range of transport ports. For example, suppose that the goal is to match TCP source ports 1000 to 1999, inclusive. One way is to insert 1000 flows, each of which matches on a single source port. Another way is to look at the binary representations of 1000 and 1999, as follows:

```
01111101000
11111001111
```

and then to transform those into a series of bitwise matches that accomplish the same results:

```
01111101xxx
0111111xxxx
10xxxxxxxxx
110xxxxxxxx
1110xxxxxxx
11110xxxxxx
1111100xxxx
```

which become the following when written in the syntax required by `ovs-ofctl`:

```
tcp,tcp_src=0x03e8/0xffff8
tcp,tcp_src=0x03f0/0xffff0
tcp,tcp_src=0x0400/0xfe00
tcp,tcp_src=0x0600/0xff00
tcp,tcp_src=0x0700/0xff80
tcp,tcp_src=0x0780/0xffc0
tcp,tcp_src=0x07c0/0xffff0
```

Only Open vSwitch 1.6 and later supports bitwise matching on transport ports.

Like the exact-match forms described above, the bitwise match forms apply only when `dl_type` and `nw_proto` specify TCP or UDP or SCTP.

`tp_src=port`

`tp_dst=port`

These are deprecated generic forms of L4 port matches. In new code, please use the TCP-, UDP-, or SCTP-specific forms described above.

`tcp_flags=flags/mask`

`tcp_flags=[+flag...][-flag...]`

Bitwise match on TCP flags. The flags and mask are 16-bit numbers written in decimal or in hexadecimal prefixed by 0x. Each 1-bit in mask requires that the corresponding bit in flags must match. Each 0-bit in mask causes the corresponding bit to be ignored.

Alternatively, the flags can be specified by their symbolic names (listed below), each preceded by either + for a flag that must be set, or - for a flag that must be unset, without any other delimiters between the flags. Flags not mentioned are wildcarded. For example, `tcp,tcp_flags=+syn-ack` matches TCP SYNs that are not ACKs.

TCP protocol currently defines 9 flag bits, and additional 3 bits are reserved (must be transmitted as zero), see RFCs 793, 3168, and 3540. The flag bits are, numbering from the least significant bit:

0: fin No more data from sender.

1: syn Synchronize sequence numbers.

2: rst Reset the connection.

3: psh Push function.

4: ack Acknowledgement field significant.

5: urg Urgent pointer field significant.

6: ece ECN Echo.

7: cwr Congestion Windows Reduced.

8: ns Nonce Sum.

9-11: Reserved.

12-15: Not matchable, must be zero.

`icmp_type=type`

`icmp_code=code`

When `dl_type` and `nw_proto` specify ICMP or ICMPv6, type matches the ICMP type and code matches the ICMP code. Each is specified as a decimal number between 0 and 255, inclusive.

When `dl_type` and `nw_proto` take other values, the values of these settings are ignored (see Flow Syntax above).

`table=number`

For flow dump commands, limits the flows dumped to those in the table with the given number between 0 and 254. If not specified (or if 255 is specified as number), then flows in all tables are dumped.

For flow table modification commands, behavior varies based on the OpenFlow version used to connect to the switch:

OpenFlow 1.0

OpenFlow 1.0 does not support table for modifying flows. `ovs-ofctl` will exit with an error if table (other than `table=255`) is specified for a switch that only supports OpenFlow 1.0.

In OpenFlow 1.0, the switch chooses the table into which to insert a new flow. The Open vSwitch software switch always chooses table 0. Other Open vSwitch datapaths and other OpenFlow implementations may choose different tables.

The OpenFlow 1.0 behavior in Open vSwitch for modifying or removing flows depends on whether `--strict` is used. Without `--strict`, the command applies to matching flows in all tables. With `--strict`, the command will operate on any single matching flow in any table; it will do nothing if there are matches in more than one table. (The distinction between these behaviors only matters if non-OpenFlow 1.0 commands were also used, because OpenFlow 1.0 alone cannot add flows with the same matching criteria to multiple tables.)

OpenFlow 1.0 with `table_id` extension

Open vSwitch implements an OpenFlow extension that allows the controller to specify the table on which to operate. `ovs-ofctl` automatically enables the extension when table is specified and OpenFlow 1.0 is used. `ovs-ofctl` automatically detects whether the switch supports the extension. As of this writing, this extension is only known to be implemented by Open vSwitch.

With this extension, `ovs-ofctl` operates on the requested table when table is specified, and acts as described for OpenFlow 1.0 above when no table is specified (or for `table=255`).

OpenFlow 1.1

OpenFlow 1.1 requires flow table modification commands to specify a table. When table is not specified (or `table=255` is specified), `ovs-ofctl` defaults to table 0.

OpenFlow 1.2 and later

OpenFlow 1.2 and later allow flow deletion commands, but not other flow table modification commands, to operate on all flow tables, with the behavior described above for OpenFlow 1.0.

`metadata=value[/mask]`

Matches value either exactly or with optional mask in the metadata field. value and mask are 64-bit integers, by default in decimal (use a 0x prefix to specify hexadecimal). Arbitrary mask values are allowed: a 1-bit in mask indicates that the corresponding bit in value must match exactly, and a 0-bit wildcards that bit. Matching on metadata was added in Open vSwitch 1.8.

The following shorthand notations are also available:

```
ip      Same as dl_type=0x0800.
ipv6    Same as dl_type=0x86dd.
icmp    Same as dl_type=0x0800,nw_proto=1.
icmp6   Same as dl_type=0x86dd,nw_proto=58.
tcp     Same as dl_type=0x0800,nw_proto=6.
tcp6    Same as dl_type=0x86dd,nw_proto=6.
udp     Same as dl_type=0x0800,nw_proto=17.
udp6    Same as dl_type=0x86dd,nw_proto=17.
sctp    Same as dl_type=0x0800,nw_proto=132.
sctp6   Same as dl_type=0x86dd,nw_proto=132.
arp     Same as dl_type=0x0806.
rarp    Same as dl_type=0x8035.
mpls    Same as dl_type=0x8847.
mplsm   Same as dl_type=0x8848.
```

The following field assignments require support for the NXM (Nicira Extended Match) extension to OpenFlow. When one of these is specified, ovs-ofctl will automatically attempt to negotiate use of this extension. If the switch does not support NXM, then ovs-ofctl will report a fatal error.

`vlan_tci=tci[/mask]`

Matches modified VLAN TCI tci. If mask is omitted, tci is the exact VLAN TCI to match; if mask is specified, then a 1-bit in mask indicates that the corresponding bit in tci must match exactly, and a 0-bit wildcards that bit. Both tci and mask are 16-bit values that are decimal by default; use a 0x prefix to specify them in hexadecimal.

The value that vlan_tci matches against is 0 for a packet that has no 802.1Q header. Otherwise, it is the TCI value from the 802.1Q header with the CFI bit (with value 0x1000) forced to 1.

Examples:

```
vlan_tci=0
    Match packets without an 802.1Q header.
```

`vlan_tci=0x1000/0x1000`
 Match packets with an 802.1Q header, regardless of VLAN and priority values.

`vlan_tci=0xf123`
 Match packets tagged with priority 7 in VLAN 0x123.

`vlan_tci=0x1123/0x1fff`
 Match packets tagged with VLAN 0x123 (and any priority).

`vlan_tci=0x5000/0xf000`
 Match packets tagged with priority 2 (in any VLAN).

`vlan_tci=0/0xffff`
 Match packets with no 802.1Q header or tagged with VLAN 0 (and any priority).

`vlan_tci=0x5000/0xe000`
 Match packets with no 802.1Q header or tagged with priority 2 (in any VLAN).

`vlan_tci=0/0xefff`
 Match packets with no 802.1Q header or tagged with VLAN 0 and priority 0.

Some of these matching possibilities can also be achieved with `dl_vlan` and `dl_vlan_pcp`.

`ip_frag=frag_type`
 When `dl_type` specifies IP or IPv6, `frag_type` specifies what kind of IP fragments or non-fragments to match. The following values of `frag_type` are supported:

`no` Matches only non-fragmented packets.

`yes` Matches all fragments.

`first` Matches only fragments with offset 0.

`later` Matches only fragments with nonzero offset.

`not_later`
 Matches non-fragmented packets and fragments with zero offset.

The `ip_frag` match type is likely to be most useful in `nx-match` mode. See the description of the `set-frags` command, above, for more details.

`arp_spa=ip[/netmask]`

`arp_tpa=ip[/netmask]`

When `dl_type` specifies either ARP or RARP, `arp_spa` and `arp_tpa` match the source and target IPv4 address, respectively. An address may be specified as an IP address or host name (e.g. 192.168.1.1 or `www.example.com`). The optional netmask allows restricting a match to an IPv4 address prefix. The netmask may be specified as a dotted quad (e.g. 192.168.1.0/255.255.255.0) or as a CIDR block (e.g. 192.168.1.0/24).

`arp_sha=xx:xx:xx:xx:xx:xx`

`arp_tha=xx:xx:xx:xx:xx:xx`

When `dl_type` specifies either ARP or RARP, `arp_sha` and `arp_tha` match the source and target hardware address, respectively. An address is specified as 6 pairs of hexadecimal digits delimited by colons (e.g. 00:0A:E4:25:6B:B0).

`arp_sha=xx:xx:xx:xx:xx:xx/xx:xx:xx:xx:xx:xx`

`arp_tha=xx:xx:xx:xx:xx:xx/xx:xx:xx:xx:xx:xx`

When `dl_type` specifies either ARP or RARP, `arp_sha` and `arp_tha` match the source and target hardware address, respectively. An address is specified as 6 pairs of hexadecimal digits delimited by colons (e.g. 00:0A:E4:25:6B:B0), with a wildcard mask following the slash.

`arp_op=opcode`

When `dl_type` specifies either ARP or RARP, `arp_op` matches the ARP opcode. Only ARP opcodes between 1 and 255 should be specified for matching.

`ipv6_src=ipv6[/netmask]`

`ipv6_dst=ipv6[/netmask]`

When `dl_type` is 0x86dd (possibly via shorthand, e.g., `ipv6` or `tcp6`), matches IPv6 source (or destination) address `ipv6`, which may be specified as defined in RFC 2373. The preferred format is `x:x:x:x:x:x:x:x`, where `x` are the hexadecimal values of the eight 16-bit pieces of the address. A single instance of `::` may be used to indicate multiple groups of 16-bits of zeros. The optional `netmask` allows restricting a match to an IPv6 address prefix. A `netmask` is specified as an IPv6 address (e.g. 2001:db8:3c4d:1::ffff:ffff:ffff:ffff::) or a CIDR block (e.g. 2001:db8:3c4d:1::/64). Open vSwitch 1.8 and later support arbitrary masks; earlier versions support only CIDR masks, that is, CIDR block and IPv6 addresses that are equivalent to CIDR blocks.

`ipv6_label=label`

When `dl_type` is 0x86dd (possibly via shorthand, e.g., `ipv6` or `tcp6`), matches IPv6 flow label `label`.

`nd_target=ipv6[/netmask]`

When `dl_type`, `nw_proto`, and `icmp_type` specify IPv6 Neighbor Discovery (ICMPv6 type 135 or 136), matches the target address `ipv6`. `ipv6` is in the same format described earlier for the `ipv6_src` and `ipv6_dst` fields.

`nd_sll=xx:xx:xx:xx:xx:xx`

When `dl_type`, `nw_proto`, and `icmp_type` specify IPv6 Neighbor Solicitation (ICMPv6 type 135), matches the source link-layer address option. An address is specified as 6 pairs of hexadecimal digits delimited by colons.

`nd_tll=xx:xx:xx:xx:xx:xx`

When `dl_type`, `nw_proto`, and `icmp_type` specify IPv6 Neighbor Advertisement (ICMPv6 type 136), matches the target link-layer address option. An address is specified as 6 pairs of hexadecimal digits delimited by colons.

`mpls_bos=bos`

When `dl_type` is 0x8847 or 0x8848 (possibly via shorthand e.g., `mpls` or `mplsm`), matches the bottom-of-stack bit of the outer-

most MPLS label stack entry. Valid values are 0 and 1.

If 1 then for a packet with a well-formed MPLS label stack the bottom-of-stack bit indicates that the outer label stack entry is also the inner-most label stack entry and thus that is that there is only one label stack entry present. Conversely, if 0 then for a packet with a well-formed MPLS label stack the bottom-of-stack bit indicates that the outer label stack entry is not the inner-most label stack entry and thus there is more than one label stack entry present.

`mpls_label=label`

When `dl_type` is 0x8847 or 0x8848 (possibly via shorthand e.g., `mpls` or `mplsm`), matches the label of the outer MPLS label stack entry. The label is a 20-bit value that is decimal by default; use a 0x prefix to specify them in hexadecimal.

`mpls_tc=tc`

When `dl_type` is 0x8847 or 0x8848 (possibly via shorthand e.g., `mpls` or `mplsm`), matches the traffic-class of the outer MPLS label stack entry. Valid values are between 0 (lowest) and 7 (highest).

`regidx=value[/mask]`

Matches `value` either exactly or with optional mask in register number `idx`. The valid range of `idx` depends on the switch. `value` and `mask` are 32-bit integers, by default in decimal (use a 0x prefix to specify hexadecimal). Arbitrary mask values are allowed: a 1-bit in mask indicates that the corresponding bit in `value` must match exactly, and a 0-bit wildcards that bit.

When a packet enters an OpenFlow switch, all of the registers are set to 0. Only explicit actions change register values.

`xregidx=value[/mask]`

Matches `value` either exactly or with optional mask in 64-bit ``extended register'' number `idx`. Each of the 64-bit extended registers overlays two of the 32-bit registers: `xreg0` overlays `reg0` and `reg1`, with `reg0` supplying the most-significant bits of `xreg0` and `reg1` the least-significant. `xreg1` similarly overlays `reg2` and `reg3`, and so on.

These fields were added in Open vSwitch 2.3 to conform with the OpenFlow 1.5 specification. OpenFlow 1.5 calls these fields just the ``packet registers,'' but Open vSwitch already had 32-bit registers by that name, which is why Open vSwitch refers to the standard registers as ``extended registers''.

`pkt_mark=value[/mask]`

Matches packet metadata mark value either exactly or with optional mask. The mark is associated data that may be passed into other system components in order to facilitate interaction between subsystems. On Linux this corresponds to the `skb mark` but the exact implementation is platform-dependent.

`actset_output=port`

Matches the output port currently in the OpenFlow action set, where `port` may be an OpenFlow port number or keyword (e.g. `LOCAL`). If there is no output port in the OpenFlow action set, or if the output port will be ignored (e.g. because there is an

output group in the OpenFlow action set), then the value will be UNSET.

This field was introduced in Open vSwitch 2.4 to conform with the OpenFlow 1.5 specification.

`conj_id=value`

Matches the given 32-bit value against the conjunction ID. This is used only with the conjunction action (see below).

This field was introduced in Open vSwitch 2.4.

`ct_state=flags/mask`

`ct_state=[+flag...][-flag...]`

Bitwise match on connection state flags. This is used with the ct action (see below).

The `ct_state` field provides information from a connection tracking module. It describes whether the packet has previously traversed the connection tracker (tracked, or `trk`) and, if it has been tracked, any additional information that the connection tracker was able to provide about the connection that the current packet belongs to.

Individual packets may be in one of two states: Untracked or tracked. When the ct action is executed on a packet, it becomes tracked for the remainder of OpenFlow pipeline processing. Once a packet has become tracked, the state of its corresponding connection may be determined. Note that the `ct_state` is only significant for the current `ct_zone`.

Connections may be in one of two states: uncommitted or committed. Connections are uncommitted by default. To determine ongoing information about a connection, like whether the connection is established or not, the connection must be committed. When the ct action is executed on a packet with the `commit` parameter, the connection will become committed and will remain in this state until the end of the connection. Committed connections store state beyond the duration of packet processing.

The `flags` and `mask` are 32-bit numbers written in decimal or in hexadecimal prefixed by `0x`. Each 1-bit in `mask` requires that the corresponding bit in `flags` must match. Each 0-bit in `mask` causes the corresponding bit to be ignored.

Alternatively, the flags can be specified by their symbolic names (listed below), each preceded by either `+` for a flag that must be set, or `-` for a flag that must be unset, without any other delimiters between the flags. Flags not mentioned are wildcarded. For example, `tcp,ct_state=+trk-new` matches TCP packets that have been run through the connection tracker and do not establish a new connection.

The following flags describe the state of the tracking:

`0x01: new`

This is the beginning of a new connection. This flag may only be present for uncommitted connections.

`0x02: est`

This is part of an already existing connection. This flag

may only be present for committed connections.

0x04: rel

This is a connection that is related to an existing connection, for instance ICMP "destination unreachable" messages or FTP data connections. This flag may only be present for committed connections.

0x08: rpl

The flow is in the reply direction, meaning it did not initiate the connection. This flag may only be present for committed connections.

0x10: inv

The state is invalid, meaning that the connection tracker couldn't identify the connection. This flag is a catch-all for any problems that the connection tracker may have, for example:

- L3/L4 protocol handler is not loaded/unavailable. With the Linux kernel datapath, this may mean that the "nf_conntrack_ipv4" or "nf_conntrack_ipv6" modules are not loaded.

- L3/L4 protocol handler determines that the packet is malformed.

- Packets are unexpected length for protocol.

0x20: trk

This packet is tracked, meaning that it has previously traversed the connection tracker. If this flag is not set, then no other flags will be set. If this flag is set, then the packet is tracked and other flags may also be set.

0x40: snat

This packet was transformed by source address/port translation by a preceding ct action.

0x80: dnat

This packet was transformed by destination address/port translation by a preceding ct action.

This field was introduced in Open vSwitch 2.5. The snat and dnat bits were added in Open vSwitch 2.6.

The following fields are associated with the connection tracker and will only be populated for tracked packets. The ct action will populate these fields, and allows modification of some of the below fields.

ct_zone=zone

Matches the given 16-bit connection zone exactly. This represents the most recent connection tracking context that ct was executed in. Each zone is an independent connection tracking context, so if you wish to track the same packet in multiple contexts then you must use the ct action multiple times. Introduced in Open vSwitch 2.5.

ct_mark=value[/mask]

Matches the given 32-bit connection mark value either exactly or

with optional mask. This represents metadata associated with the connection that the current packet is part of. Introduced in Open vSwitch 2.5.

`ct_label=value[/mask]`

Matches the given 128-bit connection labels value either exactly or with optional mask. This represents metadata associated with the connection that the current packet is part of. Introduced in Open vSwitch 2.5.

Defining IPv6 flows (those with `dl_type` equal to 0x86dd) requires support for NXM. The following shorthand notations are available for IPv6-related flows:

`ipv6` Same as `dl_type=0x86dd`.

`tcp6` Same as `dl_type=0x86dd,nw_proto=6`.

`udp6` Same as `dl_type=0x86dd,nw_proto=17`.

`sctp6` Same as `dl_type=0x86dd,nw_proto=132`.

`icmp6` Same as `dl_type=0x86dd,nw_proto=58`.

Finally, field assignments to `duration`, `n_packets`, or `n_bytes` are ignored to allow output from the `dump-flows` command to be used as input for other commands that parse flows.

The `add-flow`, `add-flows`, and `mod-flows` commands require an additional field, which must be the final field specified:

`actions=[action][,action...]`

Specifies a comma-separated list of actions to take on a packet when the flow entry matches. If no action is specified, then packets matching the flow are dropped. The following forms of action are supported:

`port`

`output:port`

Outputs the packet to OpenFlow port number `port`. If `port` is the packet's input port, the packet is not output.

`output:src[start..end]`

Outputs the packet to the OpenFlow port number read from `src`, which must be an NXM field as described above. For example, `output:NXM_NX_REG0[16..31]` outputs to the OpenFlow port number written in the upper half of register 0. If the port number is the packet's input port, the packet is not output.

This form of output was added in Open vSwitch 1.3.0. This form of output uses an OpenFlow extension that is not supported by standard OpenFlow switches.

`output(port=port,max_len=nbytes)`

Outputs the packet to the OpenFlow port number read from `port`, with maximum packet size set to `nbytes`. `port` may be OpenFlow port number, `local`, or `in_port`. Patch `port` is not supported. Packets larger than `nbytes` will be trimmed to `nbytes` while packets smaller than `nbytes` remains the original size.

`group:group_id`
Outputs the packet to the OpenFlow group `group_id`. Group tables are only supported in OpenFlow 1.1+. See Group Syntax for more details.

`normal` Subjects the packet to the device's normal L2/L3 processing. (This action is not implemented by all OpenFlow switches.)

`flood` Outputs the packet on all switch physical ports other than the port on which it was received and any ports on which flooding is disabled (typically, these would be ports disabled by the IEEE 802.1D spanning tree protocol).

`all` Outputs the packet on all switch physical ports other than the port on which it was received.

`local` Outputs the packet on the ``local port,'' which corresponds to the network device that has the same name as the bridge.

`in_port`
Outputs the packet on the port from which it was received.

`controller(key=value...)`
Sends the packet and its metadata to the OpenFlow controller as a ``packet in'' message. The supported key-value pairs are:

`max_len=nbytes`
Limit to `nbytes` the number of bytes of the packet to send to the controller. By default the entire packet is sent.

`reason=reason`
Specify `reason` as the reason for sending the message in the ``packet in'' message. The supported reasons are `action` (the default), `no_match`, and `invalid_ttl`.

`id=controller-id`
Specify `controller-id`, a 16-bit integer, as the connection ID of the OpenFlow controller or controllers to which the ``packet in'' message should be sent. The default is zero. Zero is also the default connection ID for each controller connection, and a given controller connection will only have a nonzero connection ID if its controller uses the `NXT_SET_CONTROLLER_ID` Nicira extension to OpenFlow.

`userdata=hh...`
Supplies the bytes represented as hex digits `hh` as additional data to the controller in the packet-in message. Pairs of hex digits may be separated by periods for readability.

`pause` Causes the switch to freeze the packet's trip through Open vSwitch flow tables and serializes

that state into the packet-in message as a ``continuation,'' an additional property in the `NXT_PACKET_IN2` message. The controller can later send the continuation back to the switch in an `NXT_RESUME` message, which will restart the packet's traversal from the point where it was interrupted. This permits an OpenFlow controller to interpose on a packet midway through processing in Open vSwitch.

If any reason other than action or any nonzero controller-id is supplied, Open vSwitch extension `NXAST_CONTROLLER`, supported by Open vSwitch 1.6 and later, is used. If userdata is supplied, then `NXAST_CONTROLLER2`, supported by Open vSwitch 2.6 and later, is used.

`controller`

`controller[:nbytes]`

Shorthand for `controller()` or `controller(max_len=nbytes)`, respectively.

`drop` Discards the packet, so no further processing or forwarding takes place. If a drop action is used, no other actions may be specified.

`mod_vlan_vid:vlan_vid`

Modifies the VLAN id on a packet. The VLAN tag is added or modified as necessary to match the value specified. If the VLAN tag is added, a priority of zero is used (see the `mod_vlan_pcp` action to set this).

`mod_vlan_pcp:vlan_pcp`

Modifies the VLAN priority on a packet. The VLAN tag is added or modified as necessary to match the value specified. Valid values are between 0 (lowest) and 7 (highest). If the VLAN tag is added, a vid of zero is used (see the `mod_vlan_vid` action to set this).

`strip_vlan`

Strips the VLAN tag from a packet if it is present.

`push_vlan:ethertype`

Push a new VLAN tag onto the packet. Ethertype is used as the Ethertype for the tag. Only ethertype 0x8100 should be used. (0x88a8 which the spec allows isn't supported at the moment.) A priority of zero and the tag of zero are used for the new tag.

`push_mpls:ethertype`

Changes the packet's Ethertype to ethertype, which must be either 0x8847 or 0x8848, and pushes an MPLS LSE.

If the packet does not already contain any MPLS labels then an initial label stack entry is pushed. The label stack entry's label is 2 if the packet contains IPv6 and 0 otherwise, its default traffic control value is the low 3 bits of the packet's DSCP value (0 if the packet is not IP), and its TTL is copied from the IP TTL (64 if the packet is not IP).

If the packet does already contain an MPLS label, pushes a new outermost label as a copy of the existing outermost label.

A limitation of the implementation is that processing of actions will stop if push_mpls follows another push_mpls unless there is a pop_mpls in between.

`pop_mpls:ethertype`
Strips the outermost MPLS label stack entry. Currently the implementation restricts ethertype to a non-MPLS Ethertype and thus pop_mpls should only be applied to packets with an MPLS label stack depth of one. A further limitation is that processing of actions will stop if pop_mpls follows another pop_mpls unless there is a push_mpls in between.

`mod_dl_src:mac`
Sets the source Ethernet address to mac.

`mod_dl_dst:mac`
Sets the destination Ethernet address to mac.

`mod_nw_src:ip`
Sets the IPv4 source address to ip.

`mod_nw_dst:ip`
Sets the IPv4 destination address to ip.

`mod_tp_src:port`
Sets the TCP or UDP or SCTP source port to port.

`mod_tp_dst:port`
Sets the TCP or UDP or SCTP destination port to port.

`mod_nw_tos:tos`
Sets the DSCP bits in the IPv4 ToS/DSCP or IPv6 traffic class field to tos, which must be a multiple of 4 between 0 and 255. This action does not modify the two least significant bits of the ToS field (the ECN bits).

`mod_nw_ecn:ecn`
Sets the ECN bits in the IPv4 ToS or IPv6 traffic class field to ecn, which must be a value between 0 and 3, inclusive. This action does not modify the six most significant bits of the field (the DSCP bits).

Requires OpenFlow 1.1 or later.

`mod_nw_ttl:ttl`
Sets the IPv4 TTL or IPv6 hop limit field to ttl, which is specified as a decimal number between 0 and 255, inclusive. Switch behavior when setting ttl to zero is not well specified, though.

Requires OpenFlow 1.1 or later.

The following actions are Nicira vendor extensions that, as of this writing, are only known to be implemented by Open vSwitch:

`resubmit:port`

resubmit([port],[table])

Re-searches this OpenFlow flow table (or the table whose number is specified by table) with the in_port field replaced by port (if port is specified) and executes the actions found, if any, in addition to any other actions in this flow entry.

Recursive resubmit actions are obeyed up to implementation-defined limits:

- Open vSwitch 1.0.1 and earlier did not support recursion.
- Open vSwitch 1.0.2 and 1.0.3 limited recursion to 8 levels.
- Open vSwitch 1.1 and 1.2 limited recursion to 16 levels.
- Open vSwitch 1.2 through 1.8 limited recursion to 32 levels.
- Open vSwitch 1.9 through 2.0 limited recursion to 64 levels.
- Open vSwitch 2.1 through 2.5 limited recursion to 64 levels and impose a total limit of 4,096 resubmits per flow translation (earlier versions did not impose any total limit).
- Open vSwitch 2.6 and later imposes the same limits as 2.5, with one exception: resubmit from table x to any table y > x does not count against the recursion limit.

Open vSwitch before 1.2.90 did not support table.

ct

ct([argument],[argument...])

Send the packet through the connection tracker. Refer to the ct_state documentation above for possible packet and connection states. The following arguments are supported:

commit

Commit the connection to the connection tracking module. Information about the connection will be stored beyond the lifetime of the packet in the pipeline. Some ct_state flags are only available for committed connections.

table=number

Fork pipeline processing in two. The original instance of the packet will continue processing the current actions list as an untracked packet. An additional instance of the packet will be sent to the connection tracker, which will be re-injected into the OpenFlow pipeline to resume processing in table number, with the ct_state and other ct match fields set. If the table is not

specified, then the packet which is submitted to the connection tracker is not re-injected into the OpenFlow pipeline. It is strongly recommended to specify a table later than the current table to prevent loops.

zone=value

zone=src[start..end]

A 16-bit context id that can be used to isolate connections into separate domains, allowing overlapping network addresses in different zones. If a zone is not provided, then the default is to use zone zero. The zone may be specified either as an immediate 16-bit value, or may be provided from an NXM field src. The start and end pair are inclusive, and must specify a 16-bit range within the field. This value is copied to the ct_zone match field for packets which are re-injected into the pipeline using the table option.

exec([action][,action...])

Perform actions within the context of connection tracking. This is a restricted set of actions which are in the same format as their specifications as part of a flow. Only actions which modify the ct_mark or ct_label fields are accepted within the exec action, and these fields may only be modified with this option. For example:

set_field:value[/mask]->ct_mark

Store a 32-bit metadata value with the connection. If the connection is committed, then subsequent lookups for packets in this connection will populate the ct_mark flow field when the packet is sent to the connection tracker with the table specified.

set_field:value[/mask]->ct_label

Store a 128-bit metadata value with the connection. If the connection is committed, then subsequent lookups for packets in this connection will populate the ct_label flow field when the packet is sent to the connection tracker with the table specified.

The commit parameter must be specified to use exec(...).

alg=alg

Specify application layer gateway alg to track specific connection types. Supported types include:

ftp Look for negotiation of FTP data connections. If a subsequent FTP data connection arrives which is related, the ct action will set the rel flag in the ct_state field for packets sent through ct.

The commit parameter must be specified to use

alg=alg.

When committing related connections, the `ct_mark` for that connection is inherited from the current `ct_mark` stored with the original connection (ie, the connection created by `ct(alg=...)`).

`nat[(src|dst)=addr1[-addr2][:port1[-port2]][,flags]]`
Specify address and port translation for the connection being tracked. For new connections either `src` or `dst` argument must be provided to set up either source address/port translation (SNAT) or destination address/port translation (DNAT), respectively. Setting up address translation for a new connection takes effect only if the commit flag is also provided for the enclosing `ct` action. A bare `nat` action will only translate the packet being processed in the way the connection has been set up with an earlier `ct` action. Also a `nat` action with `src` or `dst`, when applied to a packet belonging to an established (rather than new) connection, will behave the same as a bare `nat`.

`src` and `dst` options take the following arguments:

`addr1[-addr2]`

The address range from which the translated address should be selected. If only one address is given, then that address will always be selected, otherwise the address selection can be informed by the optional persistent flag as described below. Either IPv4 or IPv6 addresses can be provided, but both addresses must be of the same type, and the datapath behavior is undefined in case of providing IPv4 address range for an IPv6 packet, or IPv6 address range for an IPv4 packet. IPv6 addresses must be bracketed with '[' and ']' if a port range is also given.

`port1[-port2]`

The port range from which the translated port should be selected. If only one port number is provided, then that should be selected. In case of a mapping conflict the datapath may choose any other non-conflicting port number instead, even when no port range is specified. The port number selection can be informed by the optional random and hash flags as described below.

The optional flags are:

`random` The selection of the port from the given range should be done using a fresh random number. This flag is mutually exclusive with `hash`.

`hash` The selection of the port from the given range should be done using a datapath spe-

cific hash of the packet's IP addresses and the other, non-mapped port number. This flag is mutually exclusive with random.

`persistent`

The selection of the IP address from the given range should be done so that the same mapping can be provided after the system restarts.

If an `alg` is specified for the committing `ct` action that also includes `nat` with a `src` or `dst` attribute, then the datapath tries to set up the helper to be NAT aware. This functionality is datapath specific and may not be supported by all datapaths.

`nat` was introduced in Open vSwitch 2.6. The first datapath that implements `ct nat` support is the one that ships with Linux 4.6.

The `ct` action may be used as a primitive to construct stateful firewalls by selectively committing some traffic, then matching the `ct_state` to allow established connections while denying new connections. The following flows provide an example of how to implement a simple firewall that allows new connections from port 1 to port 2, and only allows established connections to send traffic from port 2 to port 1:

```
table=0,priority=1,action=drop
table=0,priority=10,arp,action=normal
table=0,priority=100,ip,ct_state=-trk,action=ct(table=1)
table=1,in_port=1,ip,ct_state=+trk+new,action=ct(commit),2
table=1,in_port=1,ip,ct_state=+trk+est,action=2
table=1,in_port=2,ip,ct_state=+trk+new,action=drop
table=1,in_port=2,ip,ct_state=+trk+est,action=1
```

If `ct` is executed on IP (or IPv6) fragments, then the message is implicitly reassembled before sending to the connection tracker and refragmented upon output, to the original maximum received fragment size. Reassembly occurs within the context of the zone, meaning that IP fragments in different zones are not assembled together. Pipeline processing for the initial fragments is halted; When the final fragment is received, the message is assembled and pipeline processing will continue for that flow. Because packet ordering is not guaranteed by IP protocols, it is not possible to determine which IP fragment will cause message reassembly (and therefore continue pipeline processing). As such, it is strongly recommended that multiple flows should not execute `ct` to reassemble fragments from the same IP message.

Currently, connection tracking is only available on Linux kernels with the `nf_conntrack` module loaded. The `ct` action was introduced in Open vSwitch 2.5.

```
dec_ttl
dec_ttl(id1[,id2]...)
```

Decrement TTL of IPv4 packet or hop limit of IPv6 packet. If the TTL or hop limit is initially zero or decrementing would make it so, no decrement occurs, as packets reaching TTL zero must be rejected. Instead, a ``packet-in'' message with reason code OFPR_INVALID_TTL is sent to each connected controller that has enabled receiving them, if any. Processing the current set of actions then stops. However, if the current set of actions was reached through ``resubmit'' then remaining actions in outer levels resume processing.

This action also optionally supports the ability to specify a list of valid controller ids. Each of the controllers in the list will receive the ``packet_in'' message only if they have registered to receive the invalid ttl packets. If controller ids are not specified, the ``packet_in'' message will be sent only to the controllers having controller id zero which have registered for the invalid ttl packets.

note:[hh]...

Does nothing at all. Any number of bytes represented as hex digits hh may be included. Pairs of hex digits may be separated by periods for readability. The note action's format doesn't include an exact length for its payload, so the provided bytes will be padded on the right by enough bytes with value 0 to make the total number 6 more than a multiple of 8.

move:src[start..end]->dst[start..end]

Copies the named bits from field src to field dst. src and dst must be NXM field names as defined in nicira-ext.h, e.g. NXM_OF_UDP_SRC or NXM_NX_REG0. Each start and end pair, which are inclusive, must specify the same number of bits and must fit within its respective field. Shorthands for [start..end] exist: use [bit] to specify a single bit or [] to specify an entire field.

Examples: move:NXM_NX_REG0[0..5]->NXM_NX_REG1[26..31]
copies the six bits numbered 0 through 5, inclusive, in register 0 into bits 26 through 31, inclusive;
move:NXM_NX_REG0[0..15]->NXM_OF_VLAN_TCI[] copies the least significant 16 bits of register 0 into the VLAN TCI field.

In OpenFlow 1.0 through 1.4, move ordinarily uses an Open vSwitch extension to OpenFlow. In OpenFlow 1.5, move uses the OpenFlow 1.5 standard copy_field action. The ONF has also made copy_field available as an extension to OpenFlow 1.3. Open vSwitch 2.4 and later understands this extension and uses it if a controller uses it, but for backward compatibility with older versions of Open vSwitch, ovs-ofctl does not use it.

set_field:value[/mask]->dst

load:value->dst[start..end]

Loads a literal value into a field or part of a field. With set_field, value and the optional mask are given in the customary syntax for field dst, which is expressed as

a field name. For example, `set_field:00:11:22:33:44:55->eth_src` sets the Ethernet source address to 00:11:22:33:44:55. With `load`, value must be an integer value (in decimal or prefixed by 0x for hexadecimal) and `dst` is the NXM or OXM name for the field. For example, `load:0x001122334455->OXM_OF_ETH_DST[]` has the same effect as the prior `set_field` example.

The two forms exist for historical reasons. Open vSwitch 1.1 introduced `NXAST_REG_LOAD` as a Nicira extension to OpenFlow 1.0 and used `load` to express it. Later, OpenFlow 1.2 introduced a standard `OFPAT_SET_FIELD` action that was restricted to loading entire fields, so Open vSwitch added the form `set_field` with this restriction. OpenFlow 1.5 extended `OFPAT_SET_FIELD` to the point that it became a superset of `NXAST_REG_LOAD`. Open vSwitch translates either syntax as necessary for the OpenFlow version in use: in OpenFlow 1.0 and 1.1, `NXAST_REG_LOAD`; in OpenFlow 1.2, 1.3, and 1.4, `NXAST_REG_LOAD` for `load` or for loading a subfield, `OFPAT_SET_FIELD` otherwise; and OpenFlow 1.5 and later, `OFPAT_SET_FIELD`.

`push:src[start..end]`

Pushes start to end bits inclusive, in fields on top of the stack.

Example: `push:NXM_NX_REG2[0..5]` push the value stored in register 2 bits 0 through 5, inclusive, on to the internal stack.

`pop:dst[start..end]`

Pops from the top of the stack, retrieves the start to end bits inclusive, from the value popped and store them into the corresponding bits in `dst`.

Example: `pop:NXM_NX_REG2[0..5]` pops the value from top of the stack. Set register 2 bits 0 through 5, inclusive, based on bits 0 through 5 from the value just popped.

`multipath(fields, basis, algorithm, n_links, arg, dst[start..end])`

Hashes fields using basis as a universal hash parameter, then the applies multipath link selection algorithm (with parameter `arg`) to choose one of `n_links` output links numbered 0 through `n_links` minus 1, and stores the link into `dst[start..end]`, which must be an NXM field as described above.

fields must be one of the following:

`eth_src`

Hashes Ethernet source address only.

`symmetric_l4`

Hashes Ethernet source, destination, and type, VLAN ID, IPv4/IPv6 source, destination, and protocol, and TCP or SCTP (but not UDP) ports. The hash is computed so that pairs of corresponding flows in each direction hash to the same value, in environments where L2 paths are the same in each

direction. UDP ports are not included in the hash to support protocols such as VXLAN that use asymmetric ports in each direction.

`symmetric_l3l4`

Hashes IPv4/IPv6 source, destination, and protocol, and TCP or SCTP (but not UDP) ports. Like `symmetric_l4`, this is a symmetric hash, but by excluding L2 headers it is more effective in environments with asymmetric L2 paths (e.g. paths involving VRRP IP addresses on a router). Not an effective hash function for protocols other than IPv4 and IPv6, which hash to a constant zero.

`symmetric_l3l4+udp`

Like `symmetric_l3l4+udp`, but UDP ports are included in the hash. This is a more effective hash when asymmetric UDP protocols such as VXLAN are not a consideration.

`algorithm` must be one of `modulo_n`, `hash_threshold`, `hrw`, and `iter_hash`. Only the `iter_hash` algorithm uses `arg`.

Refer to `nicira-ext.h` for more details.

`bundle(fields, basis, algorithm, slave_type, slaves:[s1, s2, ...])`

Hashes `fields` using `basis` as a universal hash parameter, then applies the bundle link selection algorithm to choose one of the listed slaves represented as `slave_type`. Currently the only supported `slave_type` is `ofport`. Thus, each `s1` through `sN` should be an OpenFlow port number. Outputs to the selected slave.

Currently, `fields` must be either `eth_src`, `symmetric_l4`, `symmetric_l3l4`, or `symmetric_l3l4+udp`, and `algorithm` must be one of `hrw` and `active_backup`.

Example: `bundle(eth_src,0,hrw,ofport,slaves:4,8)` uses an Ethernet source hash with basis 0, to select between OpenFlow ports 4 and 8 using the Highest Random Weight algorithm.

Refer to `nicira-ext.h` for more details.

`bundle_load(fields, basis, algorithm, slave_type, dst[start..end], slaves:[s1, s2, ...])`

Has the same behavior as the `bundle` action, with one exception. Instead of outputting to the selected slave, it writes its selection to `dst[start..end]`, which must be an NXM field as described above.

Example: `bundle_load(eth_src, 0, hrw, ofport, NXM_NX_REG0[], slaves:4, 8)` uses an Ethernet source hash with basis 0, to select between OpenFlow ports 4 and 8 using the Highest Random Weight algorithm, and writes the selection to `NXM_NX_REG0[]`.

Refer to `nicira-ext.h` for more details.

`learn(argument[,argument]...)`

This action adds or modifies a flow in an OpenFlow table, similar to `ovs-ofctl --strict mod-flows`. The arguments specify the flow's match fields, actions, and other properties, as follows. At least one match criterion and one action argument should ordinarily be specified.

```
idle_timeout=seconds
hard_timeout=seconds
priority=value
cookie=value
send_flow_rem
```

These arguments have the same meaning as in the usual `ovs-ofctl` flow syntax.

```
fin_idle_timeout=seconds
fin_hard_timeout=seconds
```

Adds a `fin_timeout` action with the specified arguments to the new flow. This feature was added in Open vSwitch 1.5.90.

```
table=number
```

The table in which the new flow should be inserted. Specify a decimal number between 0 and 254. The default, if table is unspecified, is table 1.

```
delete_learned
```

This flag enables deletion of the learned flows when the flow with the learn action is removed. Specifically, when the last learn action with this flag and particular table and cookie values is removed, the switch deletes all of the flows in the specified table with the specified cookie.

This flag was added in Open vSwitch 2.4.

```
field=value
field[start..end]=src[start..end]
field[start..end]
```

Adds a match criterion to the new flow.

The first form specifies that field must match the literal value, e.g. `dl_type=0x0800`. All of the fields and values for `ovs-ofctl` flow syntax are available with their usual meanings.

The second form specifies that `field[start..end]` in the new flow must match `src[start..end]` taken from the flow currently being processed.

The third form is a shorthand for the second form. It specifies that `field[start..end]` in the new flow must match `field[start..end]` taken from the flow currently being processed.

```
load:value->dst[start..end]
load:src[start..end]->dst[start..end]
```

Adds a load action to the new flow.

The first form loads the literal value into bits start through end, inclusive, in field dst. Its

syntax is the same as the `load` action described earlier in this section.

The second form loads `src[start..end]`, a value from the flow currently being processed, into bits start through end, inclusive, in field `dst`.

`output:field[start..end]`

Add an `output` action to the new flow's actions, that outputs to the OpenFlow port taken from `field[start..end]`, which must be an NXM field as described above.

For best performance, segregate learned flows into a table (using `table=number`) that is not used for any other flows except possibly for a lowest-priority ``catch-all'' flow, that is, a flow with no match criteria. (This is why the default table is 1, to keep the learned flows separate from the primary flow table 0.)

`clear_actions`

Clears all the actions in the action set immediately.

`write_actions([action][,action...])`

Add the specific actions to the action set. The syntax of actions is the same as in the `actions= field`. The action set is carried between flow tables and then executed at the end of the pipeline.

The actions in the action set are applied in the following order, as required by the OpenFlow specification, regardless of the order in which they were added to the action set. Except as specified otherwise below, the action set only holds at most a single action of each type. When more than one action of a single type is written to the action set, the one written later replaces the earlier action:

1. `strip_vlan`
 `pop_mpls`
2. `push_mpls`
3. `push_vlan`
4. `dec_ttl`
 `dec_mpls_ttl`
5. `load`
 `move`
 `mod_dl_dst`
 `mod_dl_src`
 `mod_nw_dst`
 `mod_nw_src`
 `mod_nw_tos`
 `mod_nw_ecn`
 `mod_nw_ttl`
 `mod_tp_dst`
 `mod_tp_src`
 `mod_vlan_pcp`
 `mod_vlan_vid`


```
set_field
set_tunnel
set_tunnel64
```

The action set can contain any number of these actions, with cumulative effect. They will be applied in the order as added. That is, when multiple actions modify the same part of a field, the later modification takes effect, and when they modify different parts of a field (or different fields), then both modifications are applied.

```
6.    group
      output
      resubmit
```

If more than one of these actions is present, then the one listed earliest above is executed and the others are ignored, regardless of the order in which they were added to the action set. (If none of these actions is present, the action set has no real effect, because the modified packet is not sent anywhere and thus the modifications are not visible.)

Only the actions listed above may be written to the action set.

```
write_metadata:value[/mask]
```

Updates the metadata field for the flow. If mask is omitted, the metadata field is set exactly to value; if mask is specified, then a 1-bit in mask indicates that the corresponding bit in the metadata field will be replaced with the corresponding bit from value. Both value and mask are 64-bit values that are decimal by default; use a 0x prefix to specify them in hexadecimal.

```
meter:meter_id
```

Apply the meter_id before any other actions. If a meter band rate is exceeded, the packet may be dropped, or modified, depending on the meter band type. See the description of the Meter Table Commands, above, for more details.

```
goto_table:table
```

Indicates the next table in the process pipeline.

```
fin_timeout(argument[,argument])
```

This action changes the idle timeout or hard timeout, or both, of this OpenFlow rule when the rule matches a TCP packet with the FIN or RST flag. When such a packet is observed, the action reduces the rule's timeouts to those specified on the action. If the rule's existing timeout is already shorter than the one that the action specifies, then that timeout is unaffected.

argument takes the following forms:

```
idle_timeout=seconds
```

Causes the flow to expire after the given number of seconds of inactivity.

`hard_timeout=seconds`

Causes the flow to expire after the given number of seconds, regardless of activity. (seconds specifies time since the flow's creation, not since the receipt of the FIN or RST.)

This action was added in Open vSwitch 1.5.90.

`sample(argument[,argument]...)`

Samples packets and sends one sample for every sampled packet.

argument takes the following forms:

`probability=packets`

The number of sampled packets out of 65535. Must be greater or equal to 1.

`collector_set_id=id`

The unsigned 32-bit integer identifier of the set of sample collectors to send sampled packets to. Defaults to 0.

`obs_domain_id=id`

When sending samples to IPFIX collectors, the unsigned 32-bit integer Observation Domain ID sent in every IPFIX flow record. Defaults to 0.

`obs_point_id=id`

When sending samples to IPFIX collectors, the unsigned 32-bit integer Observation Point ID sent in every IPFIX flow record. Defaults to 0.

`sampling_port=port`

Sample packets on the port. It can be set as input port or output port. When this option is omitted, or specified as NONE, IPFIX does not differentiate between ingress packets and egress packets and does not export egress tunnel information. This option was added in Open vSwitch 2.5.90.

Refer to `ovs-vswitchd.conf.db(5)` for more details on configuring sample collector sets.

This action was added in Open vSwitch 1.10.90.

`exit`

This action causes Open vSwitch to immediately halt execution of further actions. Those actions which have already been executed are unaffected. Any further actions, including those which may be in other tables, or different levels of the resubmit call stack, are ignored. Actions in the action set is still executed (specify `clear_actions` before `exit` to discard them).

`conjunction(id, k/n)`

An individual OpenFlow flow can match only a single value for each field. However, situations often arise where one wants to match one of a set of values within a field or fields. For matching a single field against a set, it is straightforward and efficient to add multiple flows to

the flow table, one for each value in the set. For example, one might use the following flows to send packets with IP source address a, b, c, or d to the OpenFlow controller:

```
ip,ip_src=a actions=controller
ip,ip_src=b actions=controller
ip,ip_src=c actions=controller
ip,ip_src=d actions=controller
```

Similarly, these flows send packets with IP destination address e, f, g, or h to the OpenFlow controller:

```
ip,ip_dst=e actions=controller
ip,ip_dst=f actions=controller
ip,ip_dst=g actions=controller
ip,ip_dst=h actions=controller
```

Installing all of the above flows in a single flow table yields a disjunctive effect: a packet is sent to the controller if $ip_src \in \{a,b,c,d\}$ or $ip_dst \in \{e,f,g,h\}$ (or both). (Pedantically, if both of the above sets of flows are present in the flow table, they should have different priorities, because OpenFlow says that the results are undefined when two flows with same priority can both match a single packet.)

Suppose, on the other hand, one wishes to match conjunctively, that is, to send a packet to the controller only if both $ip_src \in \{a,b,c,d\}$ and $ip_dst \in \{e,f,g,h\}$. This requires $4 \times 4 = 16$ flows, one for each possible pairing of ip_src and ip_dst . That is acceptable for our small example, but it does not gracefully extend to larger sets or greater numbers of dimensions.

The conjunction action is a solution for conjunctive matches that is built into Open vSwitch. A conjunction action ties groups of individual OpenFlow flows into higher-level ``conjunctive flows''. Each group corresponds to one dimension, and each flow within the group matches one possible value for the dimension. A packet that matches one flow from each group matches the conjunctive flow.

To implement a conjunctive flow with conjunction, assign the conjunctive flow a 32-bit id, which must be unique within an OpenFlow table. Assign each of the $n \geq 2$ dimensions a unique number from 1 to n ; the ordering is unimportant. Add one flow to the OpenFlow flow table for each possible value of each dimension with conjunction(id, k/n) as the flow's actions, where k is the number assigned to the flow's dimension. Together, these flows specify the conjunctive flow's match condition. When the conjunctive match condition is met, Open vSwitch looks up one more flow that specifies the conjunctive flow's actions and receives its statistics. This flow is found by setting `conj_id` to the specified id and then again searching the flow table.

The following flows provide an example. Whenever the IP source is one of the values in the flows that match on the IP source (dimension 1 of 2), and the IP destination

is one of the values in the flows that match on IP destination (dimension 2 of 2), Open vSwitch searches for a flow that matches conj_id against the conjunction ID (1234), finding the first flow listed below.

```
conj_id=1234 actions=controller
ip,ip_src=10.0.0.1 actions=conjunction(1234, 1/2)
ip,ip_src=10.0.0.4 actions=conjunction(1234, 1/2)
ip,ip_src=10.0.0.6 actions=conjunction(1234, 1/2)
ip,ip_src=10.0.0.7 actions=conjunction(1234, 1/2)
ip,ip_dst=10.0.0.2 actions=conjunction(1234, 2/2)
ip,ip_dst=10.0.0.5 actions=conjunction(1234, 2/2)
ip,ip_dst=10.0.0.7 actions=conjunction(1234, 2/2)
ip,ip_dst=10.0.0.8 actions=conjunction(1234, 2/2)
```

Many subtleties exist:

- In the example above, every flow in a single dimension has the same form, that is, dimension 1 matches on ip_src, dimension 2 on ip_dst, but this is not a requirement. Different flows within a dimension may match on different bits within a field (e.g. IP network prefixes of different lengths, or TCP/UDP port ranges as bitwise matches), or even on entirely different fields (e.g. to match packets for TCP source port 80 or TCP destination port 80).
- The flows within a dimension can vary their matches across more than one field, e.g. to match only specific pairs of IP source and destination addresses or L4 port numbers.
- A flow may have multiple conjunction actions, with different id values. This is useful for multiple conjunctive flows with overlapping sets. If one conjunctive flow matches packets with both ip_src $\in \{a,b\}$ and ip_dst $\in \{d,e\}$ and a second conjunctive flow matches ip_src $\in \{b,c\}$ and ip_dst $\in \{f,g\}$, for example, then the flow that matches ip_src=b would have two conjunction actions, one for each conjunctive flow. The order of conjunction actions within a list of actions is not significant.
- A flow with conjunction actions may also include note actions for annotations, but not any other kind of actions. (They would not be useful because they would never be executed.)
- All of the flows that constitute a conjunctive flow with a given id must have the same priority. (Flows with the same id but different priorities are currently treated as different conjunctive flows, that is, currently id values need only be unique within an OpenFlow table at a given priority. This behavior isn't guaranteed to stay the same in later releases, so please use id values unique within an OpenFlow table.)
- Conjunctive flows must not overlap with each other, at a given priority, that is, any given

packet must be able to match at most one conjunctive flow at a given priority. Overlapping conjunctive flows yield unpredictable results.

- Following a conjunctive flow match, the search for the flow with `conj_id=id` is done in the same general-purpose way as other flow table searches, so one can use flows with `conj_id=id` to act differently depending on circumstances. (One exception is that the search for the `conj_id=id` flow itself ignores conjunctive flows, to avoid recursion.) If the search with `conj_id=id` fails, Open vSwitch acts as if the conjunctive flow had not matched at all, and continues searching the flow table for other matching flows.
- OpenFlow prerequisite checking occurs for the flow with `conj_id=id` in the same way as any other flow, e.g. in an OpenFlow 1.1+ context, putting a `mod_nw_src` action into the example above would require adding an ip match, like this:
`conj_id=1234,ip actions=mod_nw_src:1.2.3.4,controller`
- OpenFlow prerequisite checking also occurs for the individual flows that comprise a conjunctive match in the same way as any other flow.
- The flows that constitute a conjunctive flow do not have useful statistics. They are never updated with byte or packet counts, and so on. (For such a flow, therefore, the idle and hard timeouts work much the same way.)
- Conjunctive flows can be a useful building block for negation, that is, inequality matches like `tcp_src ≠ 80`. To implement an inequality match, convert it to a pair of range matches, e.g. $0 \leq \text{tcp_src} < 80$ and $80 < \text{tcp_src} \leq 65535$, then convert each of the range matches into a collection of bitwise matches as explained above in the description of `tcp_src`.
- Sometimes there is a choice of which flows include a particular match. For example, suppose that we added an extra constraint to our example, to match on `ip_src ∈ {a,b,c,d}` and `ip_dst ∈ {e,f,g,h}` and `tcp_dst = i`. One way to implement this is to add the new constraint to the `conj_id` flow, like this:
`conj_id=1234,tcp,tcp_dst=i
actions=mod_nw_src:1.2.3.4,controller`
but this is not recommended because of the cost of the extra flow table lookup. Instead, add the constraint to the individual flows, either in one of the dimensions or (slightly better) all of them.
- A conjunctive match must have $n \geq 2$ dimensions (otherwise a conjunctive match is not necessary). Open vSwitch enforces this.

Each dimension within a conjunctive match should ordinarily have more than one flow. Open vSwitch does not enforce this.

The conjunction action and conj_id field were introduced in Open vSwitch 2.4.

An opaque identifier called a cookie can be used as a handle to identify a set of flows:

cookie=value

A cookie can be associated with a flow using the add-flow, add-flows, and mod-flows commands. value can be any 64-bit number and need not be unique among flows. If this field is omitted, a default cookie value of 0 is used.

cookie=value/mask

When using NXM, the cookie can be used as a handle for querying, modifying, and deleting flows. value and mask may be supplied for the del-flows, mod-flows, dump-flows, and dump-aggregate commands to limit matching cookies. A 1-bit in mask indicates that the corresponding bit in cookie must match exactly, and a 0-bit wildcards that bit. A mask of -1 may be used to exactly match a cookie.

The mod-flows command can update the cookies of flows that match a cookie by specifying the cookie field twice (once with a mask for matching and once without to indicate the new value):

```
ovs-ofctl mod-flows br0 cookie=1,actions=normal
```

Change all flows' cookies to 1 and change their actions to normal.

```
ovs-ofctl mod-flows br0 cookie=1/-1,cookie=2,actions=normal
```

Update cookies with a value of 1 to 2 and change their actions to normal.

The ability to match on cookies was added in Open vSwitch 1.5.0.

The following additional field sets the priority for flows added by the add-flow and add-flows commands. For mod-flows and del-flows when --strict is specified, priority must match along with the rest of the flow specification. For mod-flows without --strict, priority is only significant if the command creates a new flow, that is, non-strict mod-flows does not match on priority and will not change the priority of existing flows. Other commands do not allow priority to be specified.

priority=value

The priority at which a wildcarded entry will match in comparison to others. value is a number between 0 and 65535, inclusive. A higher value will match before a lower one. An exact-match entry will always have priority over an entry containing wildcards, so it has an implicit priority value of 65535. When adding a flow, if the field is not specified, the flow's priority will default to 32768.

OpenFlow leaves behavior undefined when two or more flows with the same priority can match a single packet. Some users expect ``sensible'' behavior, such as more specific flows taking prece-

dence over less specific flows, but OpenFlow does not specify this and Open vSwitch does not implement it. Users should therefore take care to use priorities to ensure the behavior that they expect.

The `add-flow`, `add-flows`, and `mod-flows` commands support the following additional options. These options affect only new flows. Thus, for `add-flow` and `add-flows`, these options are always significant, but for `mod-flows` they are significant only if the command creates a new flow, that is, their values do not update or affect existing flows.

`idle_timeout=seconds`

Causes the flow to expire after the given number of seconds of inactivity. A value of 0 (the default) prevents a flow from expiring due to inactivity.

`hard_timeout=seconds`

Causes the flow to expire after the given number of seconds, regardless of activity. A value of 0 (the default) gives the flow no hard expiration deadline.

`importance=value`

Sets the importance of a flow. The flow entry eviction mechanism can use importance as a factor in deciding which flow to evict. A value of 0 (the default) makes the flow non-evictable on the basis of importance. Specify a value between 0 and 65535.

Only OpenFlow 1.4 and later support importance.

`send_flow_rem`

Marks the flow with a flag that causes the switch to generate a ``flow removed'' message and send it to interested controllers when the flow later expires or is removed.

`check_overlap`

Forces the switch to check that the flow match does not overlap that of any different flow with the same priority in the same table. (This check is expensive so it is best to avoid it.)

The `dump-flows`, `dump-aggregate`, `del-flow` and `del-flows` commands support these additional optional fields:

`out_port=port`

If set, a matching flow must include an output action to port, which must be an OpenFlow port number or name (e.g. local).

`out_group=port`

If set, a matching flow must include an group action naming group, which must be an OpenFlow group number. This field is supported in Open vSwitch 2.5 and later and requires OpenFlow 1.1 or later.

Table Entry Output

The `dump-tables` and `dump-aggregate` commands print information about the entries in a datapath's tables. Each line of output is a flow entry as described in Flow Syntax, above, plus some additional fields:

`duration=secs`

The time, in seconds, that the entry has been in the table. secs includes as much precision as the switch provides, possibly to nanosecond resolution.

`n_packets`

The number of packets that have matched the entry.

`n_bytes`

The total number of bytes from packets that have matched the entry.

The following additional fields are included only if the switch is Open vSwitch 1.6 or later and the NXM flow format is used to dump the flow (see the description of the `--flow-format` option below). The values of these additional fields are approximations only and in particular `idle_age` will sometimes become nonzero even for busy flows.

`hard_age=secs`

The integer number of seconds since the flow was added or modified. `hard_age` is displayed only if it differs from the integer part of `duration`. (This is separate from `duration` because `mod-flows` restarts the `hard_timeout` timer without zeroing `duration`.)

`idle_age=secs`

The integer number of seconds that have passed without any packets passing through the flow.

3.10 Group Syntax

Some `ovs-ofctl` commands accept an argument that describes a group or groups. Such flow descriptions comprise a series `field=value` assignments, separated by commas or white space. (Embedding spaces into a group description normally requires quoting to prevent the shell from breaking the description into multiple arguments.) Unless noted otherwise only the last instance of each field is honoured.

`group_id=id`

The integer group id of group. When this field is specified in `del-groups` or `dump-groups`, the keyword "all" may be used to designate all groups. This field is required.

`type=type`

The type of the group. The `add-group`, `add-groups` and `mod-groups` commands require this field. It is prohibited for other commands. The following keywords designate the allowed types:

`all` Execute all buckets in the group.

`select` Execute one bucket in the group. The switch should select the bucket in such a way that should implement equal load sharing is achieved. The switch may optionally select the bucket based on bucket weights.

`indirect`

Executes the one bucket in the group.

`ff`

`fast_failover`

Executes the first live bucket in the group which is associated with a live port or group.

`command_bucket_id=id`

The bucket to operate on. The insert-buckets and remove-buckets commands require this field. It is prohibited for other commands. `id` may be an integer or one of the following keywords:

- `all` Operate on all buckets in the group. Only valid when used with the remove-buckets command in which case the effect is to remove all buckets from the group.
- `first` Operate on the first bucket present in the group. In the case of the insert-buckets command the effect is to insert new buckets just before the first bucket already present in the group; or to replace the buckets of the group if there are no buckets already present in the group. In the case of the remove-buckets command the effect is to remove the first bucket of the group; or do nothing if there are no buckets present in the group.
- `last` Operate on the last bucket present in the group. In the case of the insert-buckets command the effect is to insert new buckets just after the last bucket already present in the group; or to replace the buckets of the group if there are no buckets already present in the group. In the case of the remove-buckets command the effect is to remove the last bucket of the group; or do nothing if there are no buckets present in the group.

If `id` is an integer then it should correspond to the `bucket_id` of a bucket present in the group. In case of the insert-buckets command the effect is to insert buckets just before the bucket in the group whose `bucket_id` is `id`. In case of the remove-buckets command the effect is to remove the bucket in the group whose `bucket_id` is `id`. It is an error if there is no bucket present in the group whose `bucket_id` is `id`.

`selection_method=method`

The selection method used to select a bucket for a select group. This is a string of 1 to 15 bytes in length known to lower layers. This field is optional for add-group, add-groups and mod-group commands on groups of type select. Prohibited otherwise. The default value is the empty string.

Other than the empty string, hash is currently the only defined selection method.

This option will use a Netronome OpenFlow extension which is only supported when using Open vSwitch 2.4 and later with OpenFlow 1.5 and later.

`selection_method_param=param`

64-bit integer parameter to the selection method selected by the `selection_method` field. The parameter's use is defined by the lower-layer that implements the `selection_method`. It is optional if the `selection_method` field is specified as a non-empty string. Prohibited otherwise. The default value is zero.

This option will use a Netronome OpenFlow extension which is only supported when using Open vSwitch 2.4 and later with Open-

Flow 1.5 and later.

`fields=field`

`fields(field[=mask]...)`

The field parameters to selection method selected by the `selection_method` field. The syntax is described in Flow Syntax with the additional restrictions that if a value is provided it is treated as a wildcard mask and wildcard masks following a slash are prohibited. The pre-requisites of fields must be provided by any flows that output to the group. The use of the fields is defined by the lower-layer that implements the `selection_method`. They are optional if the `selection_method` field is specified as a non-empty string. Prohibited otherwise. The default is no fields.

This option will use a Netronome OpenFlow extension which is only supported when using Open vSwitch 2.4 and later with OpenFlow 1.5 and later.

`bucket=bucket_parameters`

The `add-group`, `add-groups` and `mod-group` commands require at least one bucket field. Bucket fields must appear after all other fields. Multiple bucket fields to specify multiple buckets. The order in which buckets are specified corresponds to their order in the group. If the type of the group is "indirect" then only one group may be specified. `bucket_parameters` consists of a list of `field=value` assignments, separated by commas or white space followed by a comma-separated list of actions. The fields for `bucket_parameters` are:

`bucket_id=id`

The 32-bit integer group id of the bucket. Values greater than 0xffffffff00 are reserved. This field was added in Open vSwitch 2.4 to conform with the OpenFlow 1.5 specification. It is not supported when earlier versions of OpenFlow are used. Open vSwitch will automatically allocate bucket ids when they are not specified.

`actions=[action][,action...]`

The syntax of actions are identical to the `actions=field` described in Flow Syntax above. Specifying `actions=` is optional, any unknown bucket parameter will be interpreted as an action.

`weight=value`

The relative weight of the bucket as an integer. This may be used by the switch during bucket select for groups whose type is select.

`watch_port=port`

Port used to determine liveness of group. This or the `watch_group` field is required for groups whose type is `ff` or `fast_failover`.

`watch_group=group_id`

Group identifier of group used to determine liveness of group. This or the `watch_port` field is required for groups whose type is `ff` or `fast_failover`.

3.11 Meter Syntax

The meter table commands accept an argument that describes a meter. Such meter descriptions comprise a series field=value assignments, separated by commas or white space. (Embedding spaces into a group description normally requires quoting to prevent the shell from breaking the description into multiple arguments.). Unless noted otherwise only the last instance of each field is honored.

meter=id

The integer meter id of the meter. When this field is specified in `del-meter`, `dump-meter`, or `meter-stats`, the keyword "all" may be used to designate all meters. This field is required, except for `meter-stats`, which dumps all stats when this field is not specified.

kbps

pktps The unit for the meter band rate parameters, either kilobits per second, or packets per second, respectively. One of these must be specified. The burst size unit corresponds to the rate unit by dropping the "per second", i.e., burst is in units of kilobits or packets, respectively.

burst Specify burst size for all bands, or none of them, if this flag is not given.

stats Collect meter and band statistics.

bands=band_parameters

The `add-meter` and `mod-meter` commands require at least one band specification. Bands must appear after all other fields.

type=type

The type of the meter band. This keyword starts a new band specification. Each band specifies a rate above which the band is to take some action. The action depends on the band type. If multiple bands' rate is exceeded, then the band with the highest rate among the exceeded bands is selected. The following keywords designate the allowed meter band types:

drop Drop packets exceeding the band's rate limit.

The other band_parameters are:

rate=value

The relative rate limit for this band, in kilobits per second or packets per second, depending on the meter flags defined above.

burst_size=size

The maximum burst allowed for the band. If `pktps` is specified, then `size` is a packet count, otherwise it is in kilobits. If unspecified, the switch is free to select some reasonable value depending on its configuration.

3.12 OPTIONS

`--strict`
Uses strict matching when running flow modification commands.

`--bundle`
Execute flow mods as an OpenFlow 1.4 atomic bundle transaction.

- Within a bundle, all flow mods are processed in the order they appear and as a single atomic transaction, meaning that if one of them fails, the whole transaction fails and none of the changes are made to the switch's flow table, and that each given datapath packet traversing the OpenFlow tables sees the flow tables either as before the transaction, or after all the flow mods in the bundle have been successfully applied.
- The beginning and the end of the flow table modification commands in a bundle are delimited with OpenFlow 1.4 bundle control messages, which makes it possible to stream the included commands without explicit OpenFlow barriers, which are otherwise used after each flow table modification command. This may make large modifications execute faster as a bundle.
- Bundles require OpenFlow 1.4 or higher. An explicit `-O OpenFlow14` option is not needed, but you may need to enable OpenFlow 1.4 support for OVS by setting the OVSDB protocols column in the bridge table.

`-O [version[,version]...]`
`--protocols=[version[,version]...]`
Sets the OpenFlow protocol versions that are allowed when establishing an OpenFlow session.

The following versions are considered to be ready for general use. These protocol versions are enabled by default:

- `OpenFlow10`, for OpenFlow 1.0.

Support for the following protocol versions is provided for testing and development purposes. They are not enabled by default:

- `OpenFlow11`, for OpenFlow 1.1.
- `OpenFlow12`, for OpenFlow 1.2.
- `OpenFlow13`, for OpenFlow 1.3.

`-F format[,format...]`
`--flow-format=format[,format...]`
ovs-ofctl supports the following individual flow formats, any number of which may be listed as format:

`OpenFlow10-table_id`
This is the standard OpenFlow 1.0 flow format. All OpenFlow switches and all versions of Open vSwitch support this flow format.

`OpenFlow10+table_id`

This is the standard OpenFlow 1.0 flow format plus a Nicira extension that allows `ovs-ofctl` to specify the flow table in which a particular flow should be placed. Open vSwitch 1.2 and later supports this flow format.

`NXM-table_id` (Nicira Extended Match)

This Nicira extension to OpenFlow is flexible and extensible. It supports all of the Nicira flow extensions, such as `tun_id` and registers. Open vSwitch 1.1 and later supports this flow format.

`NXM+table_id` (Nicira Extended Match)

This combines Nicira Extended match with the ability to place a flow in a specific table. Open vSwitch 1.2 and later supports this flow format.

`OXM-OpenFlow12`

`OXM-OpenFlow13`

`OXM-OpenFlow14`

These are the standard OXM (OpenFlow Extensible Match) flow format in OpenFlow 1.2, 1.3, and 1.4, respectively.

`ovs-ofctl` also supports the following abbreviations for collections of flow formats:

`any` Any supported flow format.

`OpenFlow10`

`OpenFlow10-table_id` or `OpenFlow10+table_id`.

`NXM` `NXM-table_id` or `NXM+table_id`.

`OXM` `OXM-OpenFlow12`, `OXM-OpenFlow13`, or `OXM-OpenFlow14`.

For commands that modify the flow table, `ovs-ofctl` by default negotiates the most widely supported flow format that supports the flows being added. For commands that query the flow table, `ovs-ofctl` by default uses the most advanced format supported by the switch.

This option, where `format` is a comma-separated list of one or more of the formats listed above, limits `ovs-ofctl`'s choice of flow format. If a command cannot work as requested using one of the specified flow formats, `ovs-ofctl` will report a fatal error.

`-P format`

`--packet-in=format=format`

`ovs-ofctl` supports the following ``packet-in'` formats, in order of increasing capability:

`standard`

This uses the `OFPT_PACKET_IN` message, the standard ``packet-in'` message for any given OpenFlow version. Every OpenFlow switch that supports a given OpenFlow version supports this format.

`nxt_packet_in`

This uses the `NXT_PACKET_IN` message, which adds many of the capabilities of the OpenFlow 1.1 and later ``packet-in'` messages before those OpenFlow versions were available in Open vSwitch. Open vSwitch 1.1 and later support

this format. Only Open vSwitch 2.6 and later, however, support it for OpenFlow 1.1 and later (but there is little reason to use it with those versions of OpenFlow).

nxt_packet_in2

This uses the `NXT_PACKET_IN2` message, which is extensible and should avoid the need to define new formats later. In particular, this format supports passing arbitrary user-provided data to a controller using the `userdata` option on the controller action. Open vSwitch 2.6 and later support this format.

Without this option, `ovs-ofctl` prefers `nxt_packet_in2` if the switch supports it. Otherwise, if OpenFlow 1.0 is in use, `ovs-ofctl` prefers `nxt_packet_in` if the switch supports it. Otherwise, `ovs-ofctl` falls back to the standard `packet-in` format. When this option is specified, `ovs-ofctl` insists on the selected format. If the switch does not support the requested format, `ovs-ofctl` will report a fatal error.

Before version 2.6, Open vSwitch called standard format `openflow10` and `nxt_packet_in` format `nxm`, and `ovs-ofctl` still accepts these names as synonyms. (The name `openflow10` was a misnomer because this format actually varies from one OpenFlow version to another; it is not consistently OpenFlow 1.0 format. Similarly, when `nxt_packet_in2` was introduced, the name `nxm` became confusing because it also uses `OXM/NXM`.)

This option affects only the monitor command.

--timestamp

Print a timestamp before each received packet. This option only affects the monitor, snoop, and `ofp-parse-pcap` commands.

-m

`--more` Increases the verbosity of OpenFlow messages printed and logged by `ovs-ofctl` commands. Specify this option more than once to increase verbosity further.

--sort[=field]

--rsort[=field]

Display output sorted by flow field in ascending (`--sort`) or descending (`--rsort`) order, where `field` is any of the fields that are allowed for matching or priority to sort by priority. When `field` is omitted, the output is sorted by priority. Specify these options multiple times to sort by multiple fields.

Any given flow will not necessarily specify a value for a given field. This requires special treatment:

- A flow that does not specify any part of a field that is used for sorting is sorted after all the flows that do specify the field. For example, `--sort=tcp_src` will sort all the flows that specify a TCP source port in ascending order, followed by the flows that do not specify a TCP source port at all.
- A flow that only specifies some bits in a field is sorted as if the wildcarded bits were zero. For example, `--sort=nw_src` would sort a flow that specifies `nw_src=192.168.0.0/24` the same as `nw_src=192.168.0.0`.

These options currently affect only dump-flows output. The following options are valid on POSIX based platforms.

--pidfile[=pidfile]

Causes a file (by default, `ovs-ofctl.pid`) to be created indicating the PID of the running process. If the `pidfile` argument is not specified, or if it does not begin with `/`, then it is created in `//var/run/openvswitch`.

If `--pidfile` is not specified, no pidfile is created.

--overwrite-pidfile

By default, when `--pidfile` is specified and the specified pidfile already exists and is locked by a running process, `ovs-ofctl` refuses to start. Specify `--overwrite-pidfile` to cause it to instead overwrite the pidfile.

When `--pidfile` is not specified, this option has no effect.

--detach

Runs `ovs-ofctl` as a background process. The process forks, and in the child it starts a new session, closes the standard file descriptors (which has the side effect of disabling logging to the console), and changes its current directory to the root (unless `--no-chdir` is specified). After the child completes its initialization, the parent exits. `ovs-ofctl` detaches only when executing the monitor or snoop commands.

--monitor

Creates an additional process to monitor the `ovs-ofctl` daemon. If the daemon dies due to a signal that indicates a programming error (`SIGABRT`, `SIGALRM`, `SIGBUS`, `SIGFPE`, `SIGILL`, `SIGPIPE`, `SIGSEGV`, `SIGXCPU`, or `SIGXFSZ`) then the monitor process starts a new copy of it. If the daemon dies or exits for another reason, the monitor process exits.

This option is normally used with `--detach`, but it also functions without it.

--no-chdir

By default, when `--detach` is specified, `ovs-ofctl` changes its current working directory to the root directory after it detaches. Otherwise, invoking `ovs-ofctl` from a carelessly chosen directory would prevent the administrator from unmounting the file system that holds that directory.

Specifying `--no-chdir` suppresses this behavior, preventing `ovs-ofctl` from changing its current working directory. This may be useful for collecting core files, since it is common behavior to write core dumps into the current working directory and the root directory is not a good directory to use.

This option has no effect when `--detach` is not specified.

--no-self-confinement

By default daemon will try to self-confine itself to work with files under well-know, at build-time whitelisted directories. It is better to stick with this default behavior and not to use this flag unless some other Access Control is used to confine daemon. Note that in contrast to other access control implemen-

tations that are typically enforced from kernel-space (e.g. DAC or MAC), self-confinement is imposed from the user-space daemon itself and hence should not be considered as a full confinement strategy, but instead should be viewed as an additional layer of security.

`--user` Causes `ovs-ofctl` to run as a different user specified in "user:group", thus dropping most of the root privileges. Short forms "user" and ":group" are also allowed, with current user or group are assumed respectively. Only daemons started by the root user accepts this argument.

On Linux, daemons will be granted `CAP_IPC_LOCK` and `CAP_NET_BIND_SERVICES` before dropping root privileges. Daemons interact with datapath, such as `ovs-vswitchd`, will be granted two additional capabilities, namely `CAP_NET_ADMIN` and `CAP_NET_RAW`. The capability change will apply even if new user is "root".

On Windows, this option is not currently supported. For security reasons, specifying this option will cause the daemon process not to start.

`--unixctl=socket`
Sets the name of the control socket on which `ovs-ofctl` listens for runtime management commands (see RUNTIME MANAGEMENT COMMANDS, below). If socket does not begin with /, it is interpreted as relative to `//var/run/openvswitch`. If `--unixctl` is not used at all, the default socket is `//var/run/openvswitch/ovs-ofctl.pid.ctl`, where pid is `ovs-ofctl`'s process ID.

On Windows, uses a kernel chosen TCP port on the localhost to listen for runtime management commands. The kernel chosen TCP port value is written in a file whose absolute path is pointed by socket. If `--unixctl` is not used at all, the file is created as `ovs-ofctl.ctl` in the configured `OVS_RUNDIR` directory.

Specifying none for socket disables the control socket feature.

Public Key Infrastructure Options

`-p privkey.pem`
`--private-key=privkey.pem`
Specifies a PEM file containing the private key used as `ovs-ofctl`'s identity for outgoing SSL connections.

`-c cert.pem`
`--certificate=cert.pem`
Specifies a PEM file containing a certificate that certifies the private key specified on `-p` or `--private-key` to be trustworthy. The certificate must be signed by the certificate authority (CA) that the peer in SSL connections will use to verify it.

`-C cacert.pem`
`--ca-cert=cacert.pem`
Specifies a PEM file containing the CA certificate that `ovs-ofctl` should use to verify certificates presented to it by SSL peers. (This may be the same certificate that SSL peers use to verify the certificate specified on `-c` or `--certificate`, or it may be a different one, depending on the PKI design in use.)

`-C none`

`--ca-cert=none`
 Disables verification of certificates presented by SSL peers. This introduces a security risk, because it means that certificates cannot be verified to be those of known trusted hosts.

`-v[spec]`

`--verbose=[spec]`

Sets logging levels. Without any spec, sets the log level for every module and destination to dbg. Otherwise, spec is a list of words separated by spaces or commas or colons, up to one from each category below:

- A valid module name, as displayed by the `vlog/list` command on `ovs-appctl(8)`, limits the log level change to the specified module.
- `syslog`, `console`, or `file`, to limit the log level change to only to the system log, to the console, or to a file, respectively. (If `--detach` is specified, `ovs-ofctl` closes its standard file descriptors, so logging to the console will have no effect.)

 On Windows platform, `syslog` is accepted as a word and is only useful along with the `--syslog-target` option (the word has no effect otherwise).
- `off`, `emer`, `err`, `warn`, `info`, or `dbg`, to control the log level. Messages of the given severity or higher will be logged, and messages of lower severity will be filtered out. `off` filters out all messages. See `ovs-appctl(8)` for a definition of each log level.

Case is not significant within spec.

Regardless of the log levels set for file, logging to a file will not take place unless `--log-file` is also specified (see below).

For compatibility with older versions of OVS, `any` is accepted as a word but has no effect.

`-v`

`--verbose`

Sets the maximum logging verbosity level, equivalent to `--verbose=dbg`.

`-vPATTERN:destination:pattern`

`--verbose=PATTERN:destination:pattern`

Sets the log pattern for destination to pattern. Refer to `ovs-appctl(8)` for a description of the valid syntax for pattern.

`-vFACILITY:facility`

`--verbose=FACILITY:facility`

Sets the RFC5424 facility of the log message. facility can be one of `kern`, `user`, `mail`, `daemon`, `auth`, `syslog`, `lpr`, `news`, `uucp`, `clock`, `ftp`, `ntp`, `audit`, `alert`, `clock2`, `local0`, `local1`, `local2`, `local3`, `local4`, `local5`, `local6` or `local7`. If this option is not specified, `daemon` is used as the default for the local system `syslog` and `local0` is used while sending a message to the target provided via the `--syslog-target` option.

`--log-file[=file]`
Enables logging to a file. If file is specified, then it is used as the exact name for the log file. The default log file name used if file is omitted is `//var/log/openvswitch/ovs-ofctl.log`.

`--syslog-target=host:port`
Send syslog messages to UDP port on host, in addition to the system syslog. The host must be a numerical IP address, not a hostname.

`--syslog-method=method`
Specify method how syslog messages should be sent to syslog daemon. Following forms are supported:

- `libc`, use `libc syslog()` function. This is the default behavior. Downside of using this options is that `libc` adds fixed prefix to every message before it is actually sent to the syslog daemon over `/dev/log` UNIX domain socket.
- `unix:file`, use UNIX domain socket directly. It is possible to specify arbitrary message format with this option. However, `rsyslogd 8.9` and older versions use hard coded parser function anyway that limits UNIX domain socket use. If you want to use arbitrary message format with older `rsyslogd` versions, then use UDP socket to localhost IP address instead.
- `udp:ip:port`, use UDP socket. With this method it is possible to use arbitrary message format also with older `rsyslogd`. When sending syslog messages over UDP socket extra precaution needs to be taken into account, for example, syslog daemon needs to be configured to listen on the specified UDP port, accidental iptables rules could be interfering with local syslog traffic and there are some security considerations that apply to UDP sockets, but do not apply to UNIX domain sockets.

`--color[=when]`
Colorize the output (for some commands); when can be never, always, or auto (the default).

Only some commands support output coloring. Color names and default colors may change in future releases.

The environment variable `OVS_COLORS` can be used to specify user-defined colors and other attributes used to highlight various parts of the output. If set, its value is a colon-separated list of capabilities that defaults to `ac=01;31:dr=34:le=31:pm=36:pr=35:sp=33:vl=32`. Supported capabilities were initially designed for coloring flows from `ovs-ofctl dump-flows` switch command, and they are as follows.

`ac=01;31`
SGR substring for actions= keyword in a flow. The default is a bold red text foreground.

`dr=34` SGR substring for drop keyword. The default is a dark blue text foreground.

- le=31 SGR substring for learn= keyword in a flow. The default is a red text foreground.
- pm=36 SGR substring for flow match attribute names. The default is a cyan text foreground.
- pr=35 SGR substring for keywords in a flow that are followed by arguments inside parenthesis. The default is a magenta text foreground.
- sp=33 SGR substring for some special keywords in a flow, notably: table=, priority=, load:, output:, move:, group:, CONTROLLER:, set_field:, resubmit:, exit. The default is a yellow text foreground.
- vl=32 SGR substring for a lone flow match attribute with no field name. The default is a green text foreground.

See the Select Graphic Rendition (SGR) section in the documentation of the text terminal that is used for permitted values and their meaning as character attributes.

- h
--help Prints a brief help message to the console.
- V
--version
Prints version information to the console.

RUNTIME MANAGEMENT COMMANDS

ovs-appctl(8) can send commands to a running ovs-ofctl process. The supported commands are listed below.

exit Causes ovs-ofctl to gracefully terminate. This command applies only when executing the monitor or snoop commands.

ofctl/set-output-file file
Causes all subsequent output to go to file instead of stderr. This command applies only when executing the monitor or snoop commands.

ofctl/send ofmsg...
Sends each ofmsg, specified as a sequence of hex digits that express an OpenFlow message, on the OpenFlow connection. This command is useful only when executing the monitor command.

ofctl/barrier
Sends an OpenFlow barrier request on the OpenFlow connection and waits for a reply. This command is useful only for the monitor command.

EXAMPLES

The following examples assume that ovs-vswitchd has a bridge named br0 configured.

ovs-ofctl dump-flows br0
Prints the flow entries in the switch.

4 ovs-vsctl

NAME

ovs-vsctl - utility for querying and configuring ovs-vswitchd

SYNOPSIS

```
ovs-vsctl [options] -- [options] command [args] [-- [options] command
[args]]...
```

DESCRIPTION

The ovs-vsctl program configures ovs-vswitchd(8) by providing a high-level interface to its configuration database. See ovs-vswitchd.conf.db(5) for comprehensive documentation of the database schema.

ovs-vsctl connects to an ovsdb-server process that maintains an Open vSwitch configuration database. Using this connection, it queries and possibly applies changes to the database, depending on the supplied commands. Then, if it applied any changes, by default it waits until ovs-vswitchd has finished reconfiguring itself before it exits. (If you use ovs-vsctl when ovs-vswitchd is not running, use --no-wait.)

ovs-vsctl can perform any number of commands in a single run, implemented as a single atomic transaction against the database.

The ovs-vsctl command line begins with global options (see OPTIONS below for details). The global options are followed by one or more commands. Each command should begin with -- by itself as a command-line argument, to separate it from the following commands. (The -- before the first command is optional.) The command itself starts with command-specific options, if any, followed by the command name and any arguments. See EXAMPLES below for syntax examples.

Linux VLAN Bridging Compatibility

The ovs-vsctl program supports the model of a bridge implemented by Open vSwitch, in which a single bridge supports ports on multiple VLANs. In this model, each port on a bridge is either a trunk port that potentially passes packets tagged with 802.1Q headers that designate VLANs or it is assigned a single implicit VLAN that is never tagged with an 802.1Q header.

For compatibility with software designed for the Linux bridge, ovs-vsctl also supports a model in which traffic associated with a given 802.1Q VLAN is segregated into a separate bridge. A special form of the add-br command (see below) creates a ``fake bridge'' within an Open vSwitch bridge to simulate this behavior. When such a ``fake bridge'' is active, ovs-vsctl will treat it much like a bridge separate from its ``parent bridge,'' but the actual implementation in Open vSwitch uses only a single bridge, with ports on the fake bridge assigned the implicit VLAN of the fake bridge of which they are members. (A fake bridge for VLAN 0 receives packets that have no 802.1Q tag or a tag with VLAN 0.)

OPTIONS

The following options affect the behavior ovs-vsctl as a whole. Some individual commands also accept their own options, which are given just before the command name. If the first command on the command line has

options, then those options must be separated from the global options by --.

--db=server

Sets server as the database server that ovs-vsctl contacts to query or modify configuration. The default is unix://var/run/openvswitch/db.sock. server must take one of the following forms:

ssl:ip:port

The specified SSL port on the host at the given ip, which must be expressed as an IP address (not a DNS name) in IPv4 or IPv6 address format. If ip is an IPv6 address, then wrap ip with square brackets, e.g.: ssl:[::1]:6640. The --private-key, --certificate, and --ca-cert options are mandatory when this form is used.

tcp:ip:port

Connect to the given TCP port on ip, where ip can be IPv4 or IPv6 address. If ip is an IPv6 address, then wrap ip with square brackets, e.g.: tcp:[::1]:6640.

unix:file

On POSIX, connect to the Unix domain server socket named file.

On Windows, connect to a localhost TCP port whose value is written in file.

pssl:port[:ip]

Listen on the given SSL port for a connection. By default, connections are not bound to a particular local IP address and it listens only on IPv4 (but not IPv6) addresses, but specifying ip limits connections to those from the given ip, either IPv4 or IPv6 address. If ip is an IPv6 address, then wrap ip with square brackets, e.g.: pssl:6640:[::1]. The --private-key, --certificate, and --ca-cert options are mandatory when this form is used.

ptcp:port[:ip]

Listen on the given TCP port for a connection. By default, connections are not bound to a particular local IP address and it listens only on IPv4 (but not IPv6) addresses, but ip may be specified to listen only for connections to the given ip, either IPv4 or IPv6 address. If ip is an IPv6 address, then wrap ip with square brackets, e.g.: ptcp:6640:[::1].

punix:file

On POSIX, listen on the Unix domain server socket named file for a connection.

On Windows, listen on a kernel chosen TCP port on the localhost. The kernel chosen TCP port value is written in file.

--no-wait

Prevents ovs-vsctl from waiting for ovs-vswitchd to reconfigure itself according to the modified database. This option should be used if ovs-vswitchd is not running; otherwise, ovs-vsctl will not exit until ovs-vswitchd starts.

This option has no effect if the commands specified do not change the database.

`--no-syslog`

By default, `ovs-vsctl` logs its arguments and the details of any changes that it makes to the system log. This option disables this logging.

This option is equivalent to `--verbose=vsctl:syslog:warn`.

`--oneline`

Modifies the output format so that the output for each command is printed on a single line. New-line characters that would otherwise separate lines are printed as `\n`, and any instances of `\` that would otherwise appear in the output are doubled. Prints a blank line for each command that has no output. This option does not affect the formatting of output from the `list` or `find` commands; see Table Formatting Options below.

`--dry-run`

Prevents `ovs-vsctl` from actually modifying the database.

`-t secs`

`--timeout=secs`

By default, or with a `secs` of 0, `ovs-vsctl` waits forever for a response from the database. This option limits runtime to approximately `secs` seconds. If the timeout expires, `ovs-vsctl` will exit with a `SIGALRM` signal. (A timeout would normally happen only if the database cannot be contacted, or if the system is overloaded.)

`--retry`

Without this option, if `ovs-vsctl` connects outward to the database server (the default) then `ovs-vsctl` will try to connect once and exit with an error if the connection fails (which usually means that `ovsdb-server` is not running).

With this option, or if `--db` specifies that `ovs-vsctl` should listen for an incoming connection from the database server, then `ovs-vsctl` will wait for a connection to the database forever.

Regardless of this setting, `--timeout` always limits how long `ovs-vsctl` will wait.

Table Formatting Options

These options control the format of output from the `list` and `find` commands.

`-f format`

`--format=format`

Sets the type of table formatting. The following types of format are available:

`table` 2-D text tables with aligned columns.

`list` (default)

A list with one column per line and rows separated by a blank line.

`html` HTML tables.

```

csv      Comma-separated values as defined in RFC 4180.

json     JSON format as defined in RFC 4627. The output is a
sequence of JSON objects, each of which corresponds to
one table. Each JSON object has the following members
with the noted values:

caption
    The table's caption. This member is omitted if
    the table has no caption.

headings
    An array with one element per table column. Each
    array element is a string giving the corresponding
    column's heading.

data     An array with one element per table row. Each
element is also an array with one element per ta-
ble column. The elements of this second-level
array are the cells that constitute the table.
Cells that represent OVSDB data or data types are
expressed in the format described in the OVSDB
specification; other cells are simply expressed as
text strings.

-d format
--data=format

Sets the formatting for cells within output tables. The follow-
ing types of format are available:

string (default)
    The simple format described in the Database Values sec-
    tion below.

bare     The simple format with punctuation stripped off: [] and
{} are omitted around sets, maps, and empty columns,
items within sets and maps are space-separated, and
strings are never quoted. This format may be easier for
scripts to parse.

json     JSON.

The json output format always outputs cells in JSON format,
ignoring this option.

--no-heading
    This option suppresses the heading row that otherwise appears in
    the first row of table output.

--pretty
    By default, JSON in output is printed as compactly as possible.
    This option causes JSON in output to be printed in a more read-
    able fashion. Members of objects and elements of arrays are
    printed one per line, with indentation.

    This option does not affect JSON in tables, which is always
    printed compactly.

--bare Equivalent to --format=list --data=bare --no-headings.

```

Public Key Infrastructure Options

`-p privkey.pem`
`--private-key=privkey.pem`
Specifies a PEM file containing the private key used as ovs-vsctl's identity for outgoing SSL connections.

`-c cert.pem`
`--certificate=cert.pem`
Specifies a PEM file containing a certificate that certifies the private key specified on `-p` or `--private-key` to be trustworthy. The certificate must be signed by the certificate authority (CA) that the peer in SSL connections will use to verify it.

`-C cacert.pem`
`--ca-cert=cacert.pem`
Specifies a PEM file containing the CA certificate that ovs-vsctl should use to verify certificates presented to it by SSL peers. (This may be the same certificate that SSL peers use to verify the certificate specified on `-c` or `--certificate`, or it may be a different one, depending on the PKI design in use.)

`-C none`
`--ca-cert=none`
Disables verification of certificates presented by SSL peers. This introduces a security risk, because it means that certificates cannot be verified to be those of known trusted hosts.

`--bootstrap-ca-cert=cacert.pem`
When `cacert.pem` exists, this option has the same effect as `-C` or `--ca-cert`. If it does not exist, then ovs-vsctl will attempt to obtain the CA certificate from the SSL peer on its first SSL connection and save it to the named PEM file. If it is successful, it will immediately drop the connection and reconnect, and from then on all SSL connections must be authenticated by a certificate signed by the CA certificate thus obtained.

This option exposes the SSL connection to a man-in-the-middle attack obtaining the initial CA certificate, but it may be useful for bootstrapping.

This option is only useful if the SSL peer sends its CA certificate as part of the SSL certificate chain. The SSL protocol does not require the server to send the CA certificate.

This option is mutually exclusive with `-C` and `--ca-cert`.

`--peer-ca-cert=peer-cacert.pem`
Specifies a PEM file that contains one or more additional certificates to send to SSL peers. `peer-cacert.pem` should be the CA certificate used to sign ovs-vsctl's own certificate, that is, the certificate specified on `-c` or `--certificate`. If ovs-vsctl's certificate is self-signed, then `--certificate` and `--peer-ca-cert` should specify the same file.

This option is not useful in normal operation, because the SSL peer must already have the CA certificate for the peer to have any confidence in ovs-vsctl's identity. However, this offers a way for a new installation to bootstrap the CA certificate on

its first SSL connection.

`-v[spec]`

`--verbose=[spec]`

Sets logging levels. Without any spec, sets the log level for every module and destination to dbg. Otherwise, spec is a list of words separated by spaces or commas or colons, up to one from each category below:

- A valid module name, as displayed by the `vlog/list` command on `ovs-appctl(8)`, limits the log level change to the specified module.
- `syslog`, `console`, or `file`, to limit the log level change to only to the system log, to the console, or to a file, respectively. (If `--detach` is specified, `ovs-vsctl` closes its standard file descriptors, so logging to the console will have no effect.)

On Windows platform, `syslog` is accepted as a word and is only useful along with the `--syslog-target` option (the word has no effect otherwise).
- `off`, `emer`, `err`, `warn`, `info`, or `dbg`, to control the log level. Messages of the given severity or higher will be logged, and messages of lower severity will be filtered out. `off` filters out all messages. See `ovs-appctl(8)` for a definition of each log level.

Case is not significant within spec.

Regardless of the log levels set for file, logging to a file will not take place unless `--log-file` is also specified (see below).

For compatibility with older versions of OVS, any is accepted as a word but has no effect.

`-v`

`--verbose`

Sets the maximum logging verbosity level, equivalent to `--verbose=dbg`.

`-vPATTERN:destination:pattern`

`--verbose=PATTERN:destination:pattern`

Sets the log pattern for destination to pattern. Refer to `ovs-appctl(8)` for a description of the valid syntax for pattern.

`-vFACILITY:facility`

`--verbose=FACILITY:facility`

Sets the RFC5424 facility of the log message. facility can be one of `kern`, `user`, `mail`, `daemon`, `auth`, `syslog`, `lpr`, `news`, `uucp`, `clock`, `ftp`, `ntp`, `audit`, `alert`, `clock2`, `local0`, `local1`, `local2`, `local3`, `local4`, `local5`, `local6` or `local7`. If this option is not specified, `daemon` is used as the default for the local system `syslog` and `local0` is used while sending a message to the target provided via the `--syslog-target` option.

`--log-file[=file]`

Enables logging to a file. If file is specified, then it is used as the exact name for the log file. The default log file

name used if file is omitted is //var/log/open-switch/ovs-vsctl.log.

`--syslog-target=host:port`

Send syslog messages to UDP port on host, in addition to the system syslog. The host must be a numerical IP address, not a hostname.

`--syslog-method=method`

Specify method how syslog messages should be sent to syslog daemon. Following forms are supported:

- `libc`, use `libc syslog()` function. This is the default behavior. Downside of using this options is that `libc` adds fixed prefix to every message before it is actually sent to the syslog daemon over `/dev/log` UNIX domain socket.
- `unix:file`, use UNIX domain socket directly. It is possible to specify arbitrary message format with this option. However, `rsyslogd` 8.9 and older versions use hard coded parser function anyway that limits UNIX domain socket use. If you want to use arbitrary message format with older `rsyslogd` versions, then use UDP socket to localhost IP address instead.
- `udp:ip:port`, use UDP socket. With this method it is possible to use arbitrary message format also with older `rsyslogd`. When sending syslog messages over UDP socket extra precaution needs to be taken into account, for example, syslog daemon needs to be configured to listen on the specified UDP port, accidental iptables rules could be interfering with local syslog traffic and there are some security considerations that apply to UDP sockets, but do not apply to UNIX domain sockets.

`-h`

`--help` Prints a brief help message to the console.

`-V`

`--version`

Prints version information to the console.

COMMANDS

The commands implemented by `ovs-vsctl` are described in the sections below.

4.1 Open vSwitch Commands

These commands work with an Open vSwitch as a whole.

`init` Initializes the Open vSwitch database, if it is empty. If the database has already been initialized, this command has no effect.

Any successful `ovs-vsctl` command automatically initializes the Open vSwitch database if it is empty. This command is provided to initialize the database without executing any other command.

`show` Prints a brief overview of the database contents.

`emer-reset`

Reset the configuration into a clean state. It deconfigures OpenFlow controllers, OVSDB servers, and SSL, and deletes port mirroring, fail_mode, NetFlow, sFlow, and IPFIX configuration.

This command also removes all other-config keys from all database records, except that other-config:hwaddr is preserved if it is present in a Bridge record. Other networking configuration is left as-is.

4.2 Bridge Commands

These commands examine and manipulate Open vSwitch bridges.

`[--may-exist] add-br bridge`

Creates a new bridge named bridge. Initially the bridge will have no ports (other than bridge itself).

Without `--may-exist`, attempting to create a bridge that exists is an error. With `--may-exist`, this command does nothing if bridge already exists as a real bridge.

`[--may-exist] add-br bridge parent vlan`

Creates a ``fake bridge'' named bridge within the existing Open vSwitch bridge parent, which must already exist and must not itself be a fake bridge. The new fake bridge will be on 802.1Q VLAN vlan, which must be an integer between 0 and 4095. The parent bridge must not already have a fake bridge for vlan. Initially bridge will have no ports (other than bridge itself).

Without `--may-exist`, attempting to create a bridge that exists is an error. With `--may-exist`, this command does nothing if bridge already exists as a VLAN bridge under parent for vlan.

`[--if-exists] del-br bridge`

Deletes bridge and all of its ports. If bridge is a real bridge, this command also deletes any fake bridges that were created with bridge as parent, including all of their ports.

Without `--if-exists`, attempting to delete a bridge that does not exist is an error. With `--if-exists`, attempting to delete a bridge that does not exist has no effect.

`[--real|--fake] list-br`

Lists all existing real and fake bridges on standard output, one per line. With `--real` or `--fake`, only bridges of that type are returned.

`br-exists bridge`

Tests whether bridge exists as a real or fake bridge. If so, `ovs-vsctl` exits successfully with exit code 0. If not, `ovs-vsctl` exits unsuccessfully with exit code 2.

`br-to-vlan bridge`

If bridge is a fake bridge, prints the bridge's 802.1Q VLAN as a decimal integer. If bridge is a real bridge, prints 0.

`br-to-parent bridge`

If bridge is a fake bridge, prints the name of its parent bridge. If bridge is a real bridge, print bridge.

```
br-set-external-id bridge key [value]
```

Sets or clears an ``external ID'' value on bridge. These values are intended to identify entities external to Open vSwitch with which bridge is associated, e.g. the bridge's identifier in a virtualization management platform. The Open vSwitch database schema specifies well-known key values, but key and value are otherwise arbitrary strings.

If value is specified, then key is set to value for bridge, overwriting any previous value. If value is omitted, then key is removed from bridge's set of external IDs (if it was present).

For real bridges, the effect of this command is similar to that of a set or remove command in the external-ids column of the Bridge table. For fake bridges, it actually modifies keys with names prefixed by fake-bridge- in the Port table.

```
br-get-external-id bridge [key]
```

Queries the external IDs on bridge. If key is specified, the output is the value for that key or the empty string if key is unset. If key is omitted, the output is key=value, one per line, for each key-value pair.

For real bridges, the effect of this command is similar to that of a get command in the external-ids column of the Bridge table. For fake bridges, it queries keys with names prefixed by fake-bridge- in the Port table.

4.3 Port Commands

These commands examine and manipulate Open vSwitch ports. These commands treat a bonded port as a single entity.

```
list-ports bridge
```

Lists all of the ports within bridge on standard output, one per line. The local port bridge is not included in the list.

```
[--may-exist] add-port bridge port [column[:key]=value]...
```

Creates on bridge a new port named port from the network device of the same name.

Optional arguments set values of column in the Port record created by the command. For example, tag=9 would make the port an access port for VLAN 9. The syntax is the same as that for the set command (see Database Commands below).

Without --may-exist, attempting to create a port that exists is an error. With --may-exist, this command does nothing if port already exists on bridge and is not a bonded port.

```
[--fake-iface] add-bond bridge port iface... [column[:key]=value]...
```

Creates on bridge a new port named port that bonds together the network devices given as each iface. At least two interfaces must be named. If the interfaces are DPDK enabled then the transaction will need to include operations to explicitly set the interface type to 'dpdk'.

Optional arguments set values of column in the Port record created by the command. The syntax is the same as that for the set command (see Database Commands below).

With `--fake-iface`, a fake interface with the name `port` is created. This should only be used for compatibility with legacy software that requires it.

Without `--may-exist`, attempting to create a port that exists is an error. With `--may-exist`, this command does nothing if `port` already exists on bridge and bonds together exactly the specified interfaces.

`[--if-exists] del-port [bridge] port`

Deletes `port`. If `bridge` is omitted, `port` is removed from whatever bridge contains it; if `bridge` is specified, it must be the real or fake bridge that contains `port`.

Without `--if-exists`, attempting to delete a port that does not exist is an error. With `--if-exists`, attempting to delete a port that does not exist has no effect.

`[--if-exists] --with-iface del-port [bridge] iface`

Deletes the port named `iface` or that has an interface named `iface`. If `bridge` is omitted, the port is removed from whatever bridge contains it; if `bridge` is specified, it must be the real or fake bridge that contains the port.

Without `--if-exists`, attempting to delete the port for an interface that does not exist is an error. With `--if-exists`, attempting to delete the port for an interface that does not exist has no effect.

`port-to-br port`

Prints the name of the bridge that contains `port` on standard output.

4.4 Interface Commands

These commands examine the interfaces attached to an Open vSwitch bridge. These commands treat a bonded port as a collection of two or more interfaces, rather than as a single port.

`list-ifaces bridge`

Lists all of the interfaces within `bridge` on standard output, one per line. The local port `bridge` is not included in the list.

`iface-to-br iface`

Prints the name of the bridge that contains `iface` on standard output.

4.5 OpenFlow Controller Connectivity

`ovs-vswitchd` can perform all configured bridging and switching locally, or it can be configured to communicate with one or more external OpenFlow controllers. The switch is typically configured to connect to a primary controller that takes charge of the bridge's flow table to implement a network policy. In addition, the switch can be configured to listen to connections from service controllers. Service controllers are typically used for occasional support and maintenance, e.g. with `ovs-ofctl`.

```

get-controller bridge
    Prints the configured controller target.

del-controller bridge
    Deletes the configured controller target.

set-controller bridge target...
    Sets the configured controller target or targets.  Each target
    may use any of the following forms:

    ssl:ip[:port]
    tcp:ip[:port]
        The specified port on the host at the given ip, which
        must be expressed as an IP address (not a DNS name) in
        IPv4 or IPv6 address format.  Wrap IPv6 addresses in
        square brackets, e.g. tcp:::1:6653.  For ssl, the
        --private-key, --certificate, and --ca-cert options are
        mandatory.

        If port is not specified, it defaults to 6653.

    unix:file
        On POSIX, a Unix domain server socket named file.

        On Windows, a localhost TCP port written in file.

    pssl:[port][:ip]
    ptcp:[port][:ip]
        Listens for OpenFlow connections on port.  The default
        port is 6653.  By default, connections are allowed from
        any IPv4 address.  Specify ip as an IPv4 address or a
        bracketed IPv6 address (e.g. ptcp:6653:::1).  DNS names
        may not be used.  For pssl, the --private-key, --certifi-
        cate, and --ca-cert options are mandatory.

    punix:file
        Listens for OpenFlow connections on the Unix domain
        server socket named file.

```

4.6 Controller Failure Settings

When a controller is configured, it is, ordinarily, responsible for setting up all flows on the switch. Thus, if the connection to the controller fails, no new network connections can be set up. If the connection to the controller stays down long enough, no packets can pass through the switch at all.

If the value is standalone, or if neither of these settings is set, ovs-vswitchd will take over responsibility for setting up flows when no message has been received from the controller for three times the inactivity probe interval. In this mode, ovs-vswitchd causes the datapath to act like an ordinary MAC-learning switch. ovs-vswitchd will continue to retry connecting to the controller in the background and, when the connection succeeds, it discontinues its standalone behavior.

If this option is set to secure, ovs-vswitchd will not set up flows on its own when the controller connection fails.

```

get-fail-mode bridge
    Prints the configured failure mode.

```

```
del-fail-mode bridge
    Deletes the configured failure mode.
```

```
set-fail-mode bridge standalone|secure
    Sets the configured failure mode.
```

4.7 Manager Connectivity

These commands manipulate the `manager_options` column in the `Open_vSwitch` table and rows in the `Managers` table. When `ovsdb-server` is configured to use the `manager_options` column for OVSDDB connections (as described in `INSTALL.Linux` and in the startup scripts provided with Open vSwitch), this allows the administrator to use `ovs-vsctl` to configure database connections.

```
get-manager
    Prints the configured manager(s).
```

```
del-manager
    Deletes the configured manager(s).
```

```
set-manager target...
    Sets the configured manager target or targets. Each target may use any of the following forms:
```

```
ssl:ip:port
    The specified SSL port on the host at the given ip, which must be expressed as an IP address (not a DNS name) in IPv4 or IPv6 address format. If ip is an IPv6 address, then wrap ip with square brackets, e.g.: ssl:[::1]:6640. The --private-key, --certificate, and --ca-cert options are mandatory when this form is used.
```

```
tcp:ip:port
    Connect to the given TCP port on ip, where ip can be IPv4 or IPv6 address. If ip is an IPv6 address, then wrap ip with square brackets, e.g.: tcp:[::1]:6640.
```

```
unix:file
    On POSIX, connect to the Unix domain server socket named file.
```

On Windows, connect to a localhost TCP port whose value is written in file.

```
pssl:port[:ip]
    Listen on the given SSL port for a connection. By default, connections are not bound to a particular local IP address and it listens only on IPv4 (but not IPv6) addresses, but specifying ip limits connections to those from the given ip, either IPv4 or IPv6 address. If ip is an IPv6 address, then wrap ip with square brackets, e.g.: pssl:6640:[::1]. The --private-key, --certificate, and --ca-cert options are mandatory when this form is used.
```

```
ptcp:port[:ip]
    Listen on the given TCP port for a connection. By default, connections are not bound to a particular local IP address and it listens only on IPv4 (but not IPv6) addresses, but ip may be specified to listen only for connections to the given ip, either IPv4 or IPv6 address.
```

If ip is an IPv6 address, then wrap ip with square brackets, e.g.: `ptcp:6640:[::1]`.

`punix:file`

On POSIX, listen on the Unix domain server socket named file for a connection.

On Windows, listen on a kernel chosen TCP port on the localhost. The kernel chosen TCP port value is written in file.

4.8 SSL Configuration

When `ovs-vswitchd` is configured to connect over SSL for management or controller connectivity, the following parameters are required:

`private-key`

Specifies a PEM file containing the private key used as the virtual switch's identity for SSL connections to the controller.

`certificate`

Specifies a PEM file containing a certificate, signed by the certificate authority (CA) used by the controller and manager, that certifies the virtual switch's private key, identifying a trustworthy switch.

`ca-cert`

Specifies a PEM file containing the CA certificate used to verify that the virtual switch is connected to a trustworthy controller.

These files are read only once, at `ovs-vswitchd` startup time. If their contents change, `ovs-vswitchd` must be killed and restarted.

These SSL settings apply to all SSL connections made by the virtual switch.

`get-ssl`

Prints the SSL configuration.

`del-ssl`

Deletes the current SSL configuration.

`[--bootstrap] set-ssl private-key certificate ca-cert`

Sets the SSL configuration. The `--bootstrap` option is described below.

CA Certificate Bootstrap

Ordinarily, all of the files named in the SSL configuration must exist when `ovs-vswitchd` starts. However, if the `ca-cert` file does not exist and the `--bootstrap` option is given, then `ovs-vswitchd` will attempt to obtain the CA certificate from the controller on its first SSL connection and save it to the named PEM file. If it is successful, it will immediately drop the connection and reconnect, and from then on all SSL connections must be authenticated by a certificate signed by the CA certificate thus obtained.

This option exposes the SSL connection to a man-in-the-middle attack

obtaining the initial CA certificate, but it may be useful for bootstrapping.

This option is only useful if the controller sends its CA certificate as part of the SSL certificate chain. The SSL protocol does not require the controller to send the CA certificate.

4.9 Auto-Attach Commands

The IETF Auto-Attach SPBM draft standard describes a compact method of using IEEE 802.1AB Link Layer Discovery Protocol (LLDP) together with a IEEE 802.1aq Shortest Path Bridging (SPB) network to automatically attach network devices to individual services in a SPB network. The intent here is to allow network applications and devices using OVS to be able to easily take advantage of features offered by industry standard SPB networks. A fundamental element of the Auto-Attach feature is to map traditional VLANs onto SPB I_SIDs. These commands manage the Auto-Attach I-SID/VLAN mappings.

`add-aa-mapping bridge i-sid vlan`
Creates a new Auto-Attach mapping on bridge for i-sid and vlan.

`del-aa-mapping bridge i-sid vlan`
Deletes an Auto-Attach mapping on bridge for i-sid and vlan.

`get-aa-mapping bridge`
Lists all of the Auto-Attach mappings within bridge on standard output.

4.10 Database Commands

These commands query and modify the contents of ovsdb tables. They are a slight abstraction of the ovsdb interface and as such they operate at a lower level than other ovs-vsctl commands.

Identifying Tables, Records, and Columns

Each of these commands has a table parameter to identify a table within the database. Many of them also take a record parameter that identifies a particular record within a table. The record parameter may be the UUID for a record, and many tables offer additional ways to identify records. Some commands also take column parameters that identify a particular field within the records in a table.

The following tables are currently defined:

Open_vSwitch
Global configuration for an ovs-vswitchd. This table contains exactly one record, identified by specifying . as the record name.

Bridge Configuration for a bridge within an Open vSwitch. Records may be identified by bridge name.

Port A bridge port. Records may be identified by port name.

Interface
A network device attached to a port. Records may be identified by name.

Flow_Table

Configuration for a particular OpenFlow flow table. Records may be identified by name.

QoS Quality-of-service configuration for a Port. Records may be identified by port name.

Queue Configuration for one queue within a QoS configuration. Records may only be identified by UUID.

Mirror A port mirroring configuration attached to a bridge. Records may be identified by mirror name.

Controller

Configuration for an OpenFlow controller. A controller attached to a particular bridge may be identified by the bridge's name.

Manager

Configuration for an OVSDDB connection. Records may be identified by target (e.g. tcp:1.2.3.4).

NetFlow

A NetFlow configuration attached to a bridge. Records may be identified by bridge name.

SSL The global SSL configuration for ovs-vswitchd. The record attached to the Open_vSwitch table may be identified by specifying . as the record name.

sFlow An sFlow exporter configuration attached to a bridge. Records may be identified by bridge name.

IPFIX An IPFIX exporter configuration attached to a bridge. Records may be identified by bridge name.

Flow_Sample_Collector_Set

An IPFIX exporter configuration attached to a bridge for sampling packets on a per-flow basis using OpenFlow sample actions.

AutoAttach

Configuration for Auto Attach within a bridge.

Record names must be specified in full and with correct capitalization. Names of tables and columns are not case-sensitive, and -- and _ are treated interchangeably. Unique abbreviations are acceptable, e.g. net or n is sufficient to identify the NetFlow table.

Database Values

Each column in the database accepts a fixed type of data. The currently defined basic types, and their representations, are:

integer

A decimal integer in the range -2^{63} to $2^{63}-1$, inclusive.

real A floating-point number.

Boolean

True or false, written true or false, respectively.

string An arbitrary Unicode string, except that null bytes are not allowed. Quotes are optional for most strings that begin with an English letter or underscore and consist only of letters, underscores, hyphens, and periods. However, true and false and strings that match the syntax of UUIDs (see below) must be enclosed in double quotes to distinguish them from other basic types. When double quotes are used, the syntax is that of strings in JSON, e.g. backslashes may be used to escape special characters. The empty string must be represented as a pair of double quotes ("").

UUID Either a universally unique identifier in the style of RFC 4122, e.g. f81d4fae-7dec-11d0-a765-00a0c91e6bf6, or an @name defined by a get or create command within the same ovs-vsctl invocation.

Multiple values in a single column may be separated by spaces or a single comma. When multiple values are present, duplicates are not allowed, and order is not important. Conversely, some database columns can have an empty set of values, represented as [], and square brackets may optionally enclose other non-empty sets or single values as well.

A few database columns are ``maps'' of key-value pairs, where the key and the value are each some fixed database type. These are specified in the form key=value, where key and value follow the syntax for the column's key type and value type, respectively. When multiple pairs are present (separated by spaces or a comma), duplicate keys are not allowed, and again the order is not important. Duplicate values are allowed. An empty map is represented as {}. Curly braces may optionally enclose non-empty maps as well (but use quotes to prevent the shell from expanding other-config={0=x,1=y} into other-config=0=x other-config=1=y, which may not have the desired effect).

Database Command Syntax

`[--if-exists] [--columns=column[,column]...] list table [record]...`
 Lists the data in each specified record. If no records are specified, lists all the records in table.

If `--columns` is specified, only the requested columns are listed, in the specified order. Otherwise, all columns are listed, in alphabetical order by column name.

Without `--if-exists`, it is an error if any specified record does not exist. With `--if-exists`, the command ignores any record that does not exist, without producing any output.

`[--columns=column[,column]...] find table [column[:key]=value]...`
 Lists the data in each record in table whose column equals value or, if key is specified, whose column contains a key with the specified value. The following operators may be used where = is written in the syntax summary:

`= != < > <= >=`

Selects records in which `column[:key]` equals, does not equal, is less than, is greater than, is less than or equal to, or is greater than or equal to value, respectively.

Consider `column[:key]` and value as sets of elements. Identical sets are considered equal. Otherwise, if the

sets have different numbers of elements, then the set with more elements is considered to be larger. Otherwise, consider a element from each set pairwise, in increasing order within each set. The first pair that differs determines the result. (For a column that contains key-value pairs, first all the keys are compared, and values are considered only if the two sets contain identical keys.)

`{=}` `{!=}`

Test for set equality or inequality, respectively.

`{<=}` Selects records in which `column[:key]` is a subset of value. For example, `flood-vlans{<=}1,2` selects records in which the `flood-vlans` column is the empty set or contains 1 or 2 or both.

`{<}` Selects records in which `column[:key]` is a proper subset of value. For example, `flood-vlans{<}1,2` selects records in which the `flood-vlans` column is the empty set or contains 1 or 2 but not both.

`{>=}` `{>}`

Same as `{<=}` and `{<}`, respectively, except that the relationship is reversed. For example, `flood-vlans{>=}1,2` selects records in which the `flood-vlans` column contains both 1 and 2.

For arithmetic operators (`=` `!=` `<` `>` `<=` `>=`), when key is specified but a particular record's column does not contain key, the record is always omitted from the results. Thus, the condition `other-config:mtu!=1500` matches records that have a `mtu` key whose value is not 1500, but not those that lack an `mtu` key.

For the set operators, when key is specified but a particular record's column does not contain key, the comparison is done against an empty set. Thus, the condition `other-config:mtu{!=}1500` matches records that have a `mtu` key whose value is not 1500 and those that lack an `mtu` key.

Don't forget to escape `<` or `>` from interpretation by the shell.

If `--columns` is specified, only the requested columns are listed, in the specified order. Otherwise all columns are listed, in alphabetical order by column name.

The UUIDs shown for rows created in the same `ovs-vsctl` invocation will be wrong.

`[--if-exists] [--id=@name] get table record [column[:key]]...`

Prints the value of each specified column in the given record in table. For map columns, a key may optionally be specified, in which case the value associated with key in the column is printed, instead of the entire map.

Without `--if-exists`, it is an error if record does not exist or key is specified, if key does not exist in record. With `--if-exists`, a missing record yields no output and a missing key

prints a blank line.

If @name is specified, then the UUID for record may be referred to by that name later in the same ovs-vsctl invocation in contexts where a UUID is expected.

Both --id and the column arguments are optional, but usually at least one or the other should be specified. If both are omitted, then get has no effect except to verify that record exists in table.

--id and --if-exists cannot be used together.

[--if-exists] set table record column[:key]=value...

Sets the value of each specified column in the given record in table to value. For map columns, a key may optionally be specified, in which case the value associated with key in that column is changed (or added, if none exists), instead of the entire map.

Without --if-exists, it is an error if record does not exist. With --if-exists, this command does nothing if record does not exist.

[--if-exists] add table record column [key=]value...

Adds the specified value or key-value pair to column in record in table. If column is a map, then key is required, otherwise it is prohibited. If key already exists in a map column, then the current value is not replaced (use the set command to replace an existing value).

Without --if-exists, it is an error if record does not exist. With --if-exists, this command does nothing if record does not exist.

[--if-exists] remove table record column value...

[--if-exists] remove table record column key...

[--if-exists] remove table record column key=value...

Removes the specified values or key-value pairs from column in record in table. The first form applies to columns that are not maps: each specified value is removed from the column. The second and third forms apply to map columns: if only a key is specified, then any key-value pair with the given key is removed, regardless of its value; if a value is given then a pair is removed only if both key and value match.

It is not an error if the column does not contain the specified key or value or pair.

Without --if-exists, it is an error if record does not exist. With --if-exists, this command does nothing if record does not exist.

[--if-exists] clear table record column...

Sets each column in record in table to the empty set or empty map, as appropriate. This command applies only to columns that are allowed to be empty.

Without --if-exists, it is an error if record does not exist.

With `--if-exists`, this command does nothing if record does not exist.

`[--id=@name] create table column[:key]=value...`

Creates a new record in table and sets the initial values of each column. Columns not explicitly set will receive their default values. Outputs the UUID of the new row.

If `@name` is specified, then the UUID for the new row may be referred to by that name elsewhere in the same `ovs-vsctl` invocation in contexts where a UUID is expected. Such references may precede or follow the create command.

Caution (ovs-vsctl as example)

Records in the Open vSwitch database are significant only when they can be reached directly or indirectly from the `Open_vSwitch` table. Except for records in the `QoS` or `Queue` tables, records that are not reachable from the `Open_vSwitch` table are automatically deleted from the database. This deletion happens immediately, without waiting for additional `ovs-vsctl` commands or other database activity. Thus, a create command must generally be accompanied by additional commands within the same `ovs-vsctl` invocation to add a chain of references to the newly created record from the top-level `Open_vSwitch` record. The `EXAMPLES` section gives some examples that show how to do this.

`[--if-exists] destroy table record...`

Deletes each specified record from table. Unless `--if-exists` is specified, each records must exist.

`--all destroy table`

Deletes all records from the table.

Caution (ovs-vsctl as example)

The destroy command is only useful for records in the `QoS` or `Queue` tables. Records in other tables are automatically deleted from the database when they become unreachable from the `Open_vSwitch` table. This means that deleting the last reference to a record is sufficient for deleting the record itself. For records in these tables, destroy is silently ignored. See the `EXAMPLES` section below for more information.

`wait-until table record [column[:key]=value]...`

Waits until table contains a record named `record` whose column equals `value` or, if `key` is specified, whose column contains a key with the specified value. Any of the operators `!=`, `<`, `>`, `<=`, or `>=` may be substituted for `=` to test for inequality, less than, greater than, less than or equal to, or greater than or equal to, respectively. (Don't forget to escape `<` or `>` from interpretation by the shell.)

If no `column[:key]=value` arguments are given, this command waits only until record exists. If more than one such argument is given, the command waits until all of them are satisfied.

Caution (ovs-vsctl as example)

Usually `wait-until` should be placed at the beginning of a set of `ovs-vsctl` commands. For example, `wait-until bridge br0 -- get bridge br0 datapath_id` waits until a bridge named `br0` is created, then prints its `datapath_id` column, whereas `get bridge br0 datapath_id -- wait-until bridge br0` will abort if no bridge named `br0` exists when `ovs-vsctl` initially connects to the database.

Consider specifying `--timeout=0` along with `--wait-until`, to prevent `ovs-vsctl` from terminating after waiting only at most 5 seconds.

`comment [arg]...`

This command has no effect on behavior, but any database log record created by the command will include the command and its arguments.

EXAMPLES

Create a new bridge named `br0` and add port `eth0` to it:

```
ovs-vsctl add-br br0
ovs-vsctl add-port br0 eth0
```

Alternatively, perform both operations in a single atomic transaction:

```
ovs-vsctl add-br br0 -- add-port br0 eth0
```

Delete bridge `br0`, reporting an error if it does not exist:

```
ovs-vsctl del-br br0
```

Delete bridge `br0` if it exists:

```
ovs-vsctl --if-exists del-br br0
```

Set the `qos` column of the Port record for `eth0` to point to a new QoS record, which in turn points with its queue 0 to a new Queue record:

```
ovs-vsctl -- set port eth0 qos=@newqos -- --id=@newqos create
qos          type=linux-htb          other-config:max-rate=1000000
queues:0=@newqueue -- --id=@newqueue create queue other-con-
fig:min-rate=1000000 other-config:max-rate=1000000
```

CONFIGURATION COOKBOOK

Port Configuration

Add an `internal` port `vlan10` to bridge `br0` as a VLAN access port for VLAN 10, and configure it with an IP address:

```
ovs-vsctl add-port br0 vlan10 tag=10 -- set Interface vlan10
type=internal

ifconfig vlan10 192.168.0.123
```

Add a GRE tunnel port `gre0` to remote IP address 1.2.3.4 to bridge `br0`:

```
ovs-vsctl add-port br0 gre0 -- set Interface gre0 type=gre
options:remote_ip=1.2.3.4
```

Port Mirroring

Mirror all packets received or sent on `eth0` or `eth1` onto `eth2`, assuming that all of those ports exist on bridge `br0` (as a side-effect this

causes any packets received on eth2 to be ignored):

```
ovs-vsctl -- set Bridge br0 mirrors=@m \  
-- --id=@eth0 get Port eth0 \  
-- --id=@eth1 get Port eth1 \  
-- --id=@eth2 get Port eth2 \  
-- --id=@m create Mirror name=mymirror select-dst-  
port=@eth0,@eth1 select-src-port=@eth0,@eth1 output-port=@eth2
```

Remove the mirror created above from br0, which also destroys the Mirror record (since it is now unreferenced):

```
ovs-vsctl -- --id=@rec get Mirror mymirror \  
-- remove Bridge br0 mirrors @rec
```

The following simpler command also works:

```
ovs-vsctl clear Bridge br0 mirrors
```

Quality of Service (QoS)

Create a linux-htb QoS record that points to a few queues and use it on eth0 and eth1:

```
ovs-vsctl -- set Port eth0 qos=@newqos \  
-- set Port eth1 qos=@newqos \  
-- --id=@newqos create QoS type=linux-htb other-con-  
fig:max-rate=1000000000 queues=0=@q0,1=@q1 \  
-- --id=@q0 create Queue other-config:min-rate=100000000  
other-config:max-rate=1000000000 \  
-- --id=@q1 create Queue other-config:min-rate=500000000
```

Deconfigure the QoS record above from eth1 only:

```
ovs-vsctl clear Port eth1 qos
```

To deconfigure the QoS record from both eth0 and eth1 and then delete the QoS record (which must be done explicitly because unreferenced QoS records are not automatically destroyed):

```
ovs-vsctl -- destroy QoS eth0 -- clear Port eth0 qos -- clear  
Port eth1 qos
```

(This command will leave two unreferenced Queue records in the database. To delete them, use "ovs-vsctl list Queue" to find their UUIDs, then "ovs-vsctl destroy Queue uuid1 uuid2" to destroy each of them or use "ovs-vsctl -- --all destroy Queue" to delete all records.)

Connectivity Monitoring

Monitor connectivity to a remote maintenance point on eth0.

```
ovs-vsctl set Interface eth0 cfm_mpid=1
```


Deconfigure connectivity monitoring from above:

```
ovs-vsctl clear Interface eth0 cfm_mpid
```

NetFlow

Configure bridge br0 to send NetFlow records to UDP port 5566 on host 192.168.0.34, with an active timeout of 30 seconds:

```
ovs-vsctl -- set Bridge br0 netflow=@nf \  
-- --id=@nf create NetFlow targets="192.168.0.34:5566" \  
active-timeout=30
```

Update the NetFlow configuration created by the previous command to instead use an active timeout of 60 seconds:

```
ovs-vsctl set NetFlow br0 active_timeout=60
```

Deconfigure the NetFlow settings from br0, which also destroys the NetFlow record (since it is now unreferenced):

```
ovs-vsctl clear Bridge br0 netflow
```

sFlow

Configure bridge br0 to send sFlow records to a collector on 10.0.0.1 at port 6343, using eth1's IP address as the source, with specific sampling parameters:

```
ovs-vsctl -- --id=@s create sFlow agent=eth1 target="10.0.0.1:6343" header=128 sampling=64 polling=10 \  
-- set Bridge br0 sflow=@s
```

Deconfigure sFlow from br0, which also destroys the sFlow record (since it is now unreferenced):

```
ovs-vsctl -- clear Bridge br0 sflow
```

IPFIX

Configure bridge br0 to send one IPFIX flow record per packet sample to UDP port 4739 on host 192.168.0.34, with Observation Domain ID 123 and Observation Point ID 456, a flow cache active timeout of 1 minute (60 seconds), maximum flow cache size of 13 flows, and flows sampled on output port with tunnel info(sampling on input and output port is enabled by default if not disabled) :

```
ovs-vsctl -- set Bridge br0 ipfix=@i \  
-- --id=@i create IPFIX targets="192.168.0.34:4739" \  
obs_domain_id=123 obs_point_id=456 cache_active_timeout=60 \  
cache_max_flows=13 \  
other_config:enable-input-sampling=false other_config:enable- \  
tunnel-sampling=true
```

Deconfigure the IPFIX settings from br0, which also destroys the IPFIX record (since it is now unreferenced):

```
ovs-vsctl clear Bridge br0 ipfix
```

802.1D Spanning Tree Protocol (STP)

Configure bridge br0 to participate in an 802.1D spanning tree:

```
ovs-vsctl set Bridge br0 stp_enable=true
```

Set the bridge priority of br0 to 0x7800:

```
ovs-vsctl set Bridge br0 other_config:stp-priority=0x7800
```

Set the path cost of port eth0 to 10:

```
ovs-vsctl set Port eth0 other_config:stp-path-cost=10
```

Deconfigure STP from above:

```
ovs-vsctl set Bridge br0 stp_enable=false
```

Multicast Snooping

Configure bridge br0 to enable multicast snooping:

```
ovs-vsctl set Bridge br0 mcast_snooping_enable=true
```

Set the multicast snooping aging time br0 to 300 seconds:

```
ovs-vsctl set Bridge br0 other_config:mcast-snooping-aging-time=300
```

Set the multicast snooping table size br0 to 2048 entries:

```
ovs-vsctl set Bridge br0 other_config:mcast-snooping-table-size=2048
```

Disable flooding of unregistered multicast packets to all ports. When set to true, the switch will send unregistered multicast packets only to ports connected to multicast routers. When it is set to false, the switch will send them to all ports. This command disables the flood of unregistered packets on bridge br0.

```
ovs-vsctl set Bridge br0 other_config:mcast-snooping-disable-flood-unregistered=true
```

Enable flooding of multicast packets (except Reports) on a specific port.

```
ovs-vsctl set Port eth1 other_config:mcast-snooping-flood=true
```

Enable flooding of Reports on a specific port.

```
ovs-vsctl set Port eth1 other_config:mcast-snooping-flood-reports=true
```

Deconfigure multicasting snooping from above:

```
ovs-vsctl set Bridge br0 mcast_snooping_enable=false
```

802.1D-2004 Rapid Spanning Tree Protocol (RSTP)

Configure bridge br0 to participate in an 802.1D-2004 Rapid Spanning Tree:

```
ovs-vsctl set Bridge br0 rstp_enable=true
```

Set the bridge address of br0 to 00:aa:aa:aa:aa:aa :

```
ovs-vsctl set Bridge br0 other_config:rstp-address=00:aa:aa:aa:aa:aa
```

Set the bridge priority of br0 to 0x7000. The value must be specified in decimal notation and should be a multiple of 4096 (if not, it is rounded down to the nearest multiple of 4096). The default priority value is 0x800 (32768).

```
ovs-vsctl set Bridge br0 other_config:rstp-priority=28672
```

Set the bridge ageing time of br0 to 1000 s. The ageing time value should be between 10 s and 1000000 s. The default value is 300 s.

```
ovs-vsctl set Bridge br0 other_config:rstp-ageing-time=1000
```

Set the bridge force protocol version of br0 to 0. The force protocol version has two acceptable values: 0 (STP compatibility mode) and 2 (normal operation).

```
ovs-vsctl set Bridge br0 other_config:rstp-force-protocol-version=0
```

Set the bridge max age of br0 to 10 s. The max age value should be between 6 s and 40 s. The default value is 20 s.

```
ovs-vsctl set Bridge br0 other_config:rstp-max-age=10
```

Set the bridge forward delay of br0 to 15 s. This value should be between 4 s and 30 s. The default value is 15 s.

```
ovs-vsctl set Bridge br0 other_config:rstp-forward-delay=15
```

Set the bridge transmit hold count of br0 to 7 s. This value should be between 1 s and 10 s. The default value is 6 s.

```
ovs-vsctl set Bridge br0 other_config:rstp-transmit-hold-count=7
```

Enable RSTP on the Port eth0:

```
ovs-vsctl set Port eth0 other_config:rstp-enable=true
```

Disable RSTP on the Port eth0:

```
ovs-vsctl set Port eth0 other_config:rstp-enable=false
```

Set the priority of port eth0 to 32. The value must be specified in decimal notation and should be a multiple of 16 (if not, it is rounded down to the nearest multiple of 16). The default priority value is 0x80 (128).

```
ovs-vsctl set Port eth0 other_config:rstp-port-priority=32
```

Set the port number of port eth0 to 3:

```
ovs-vsctl set Port eth0 other_config:rstp-port-num=3
```

Set the path cost of port eth0 to 150:

```
ovs-vsctl set Port eth0 other_config:rstp-path-cost=150
```

Set the admin edge value of port eth0:

```
ovs-vsctl set Port eth0 other_config:rstp-port-admin-edge=true
```

Set the auto edge value of port eth0:

```
ovs-vsctl set Port eth0 other_config:rstp-port-auto-edge=true
```

Set the admin point to point MAC value of port eth0. Acceptable values are 0 (not point-to-point), 1 (point-to-point, the default value) or 2 (automatic detection). The auto-detection mode is not currently implemented, and the value 2 has the same effect of 0 (not point-to-point).

```
ovs-vsctl set Port eth0 other_config:rstp-admin-p2p-mac=1
```

Set the admin port state value of port eth0. true is the default value.

```
ovs-vsctl set Port eth0 other_config:rstp-admin-port-state=false
```

Set the mcheck value of port eth0:

```
ovs-vsctl set Port eth0 other_config:rstp-port-mcheck=true
```

Deconfigure RSTP from above:

```
ovs-vsctl set Bridge br0 rstp_enable=false
```

OpenFlow Version

Configure bridge br0 to support OpenFlow versions 1.0, 1.2, and 1.3:

```
ovs-vsctl set bridge br0 protocols=OpenFlow10,OpenFlow12,OpenFlow13
```

Flow Table Configuration

Limit flow table 0 on bridge br0 to a maximum of 100 flows:

```
ovs-vsctl -- --id=@ft create Flow_Table flow_limit=100 overflow_policy=refuse -- set Bridge br0 flow_tables=0=@ft
```

EXIT STATUS

- | | |
|---|--|
| 0 | Successful program execution. |
| 1 | Usage, syntax, or configuration file error. |
| 2 | The bridge argument to br-exists specified the name of a bridge that does not exist. |